Literature base

Esteban Degetau
World Bank
estebandegetau@gmail.com

Agustín Samano
World Bank
asamanopenaloza@worldbank.org

June 24, 2026

Table of contents

§1 Motivation 3
§2 Narrative identification (pillar) 7
§3 Effects of fiscal policy 11
§4 LLMs as measurement tools 15
§5 Validation (pillar) 23
§6 Generated-regressor inference (pillar) 28
§7 Datasets 30
References 32

Role: the annotated bibliography for the methodological paper — one entry per source, recording (i) its status in the library, (ii) a summary of its contents, and (iii) its relation to our project and the rest of the literature. This is the per-source knowledge base that feeds `docs/lit_review.qmd` (synthesis) and is distinct from `docs/lit_log.md` (process decisions). Maintained by the `/lit-intake` (status + stubs), `/lit-digest` (summary + relation), and `/lit-synthesize` (reads from here) skills.

The `consultation` field is the shared state: `pending` (stub, not yet read), `abstract` (read at abstract depth), or `full` (full-text read and distilled). `/lit-digest` advances it. Citekeys come from `references.bib`; sources in Zotero but not yet exported show "pending re-export" plus their Zotero item key.

Table 1: Library status by source. File = PDF in Zotero; Tag = section/role tags applied; Consult = consultation depth.

Source	Sec	Role	File	Tag	Consult
vegh_how_2015	§1	cite-only	yes	yes	abstract
jha_effectiveness_2014	§1	cite-only	yes	yes	abstract
frankel_graduation_2013	§1	cite-only	yes	yes	abstract
ilzetzki_how_2013	§1	cite-only	yes	yes	abstract
acosta_ormaechea_tax_2012	§1	cite-only	yes	yes	abstract
crispolti_cross-country_2022	§1	cite-only	yes	yes	abstract
kaminsky_when_2004	§1	cite-only	yes	yes	abstract
romer_fiscal_2019	§1	cite-only	yes	yes	abstract
romer_new_2017	§1	cite-only	yes	yes	abstract
angrist_credibility_2010	§1	cite-only	yes	yes	abstract
romer_trouble_nodate	§1	cite-only	yes	yes	abstract
burke_global_2015	§1	exemplar	yes	yes	abstract
hsiang_quantifying_2013	§1	exemplar	yes	yes	abstract
romer2010	§2	read-deep	yes	yes	full
romer	§2	read-deep	yes	yes	full
romer_presidential_2023	§2	cite-only	yes	yes	abstract
mertens2013	§2	read-deep	yes	yes	abstract
mertens_reconciliation_2014	§2	read-deep	yes	yes	abstract
cloyne_discretionary_2013	§2	read-deep	yes	yes	abstract
tillmann_disagreement_2024	§2	cite-only	yes	yes	abstract
latifi_fiscal_nodate	§2	cite-only	yes	yes	abstract
kim_understanding_2015	§2	cite-only	yes	yes	abstract
hebous_revisiting_2018	§2	cite-only	yes	yes	abstract
blanchard_empirical_nodate	§3	cite-only	yes	yes	abstract
auerbach_measuring_2012	§3	cite-only	yes	yes	abstract
gechert_what_2015	§3	cite-only	yes	yes	abstract
riera-crichton_tax_2016	§3	cite-only	yes	yes	abstract
ramey_ten_2019	§3	cite-only	yes	yes	abstract
ramey_macroeconomic_2016	§3	cite-only	yes	yes	abstract
ramey_identifying_2011	§3	cite-only	yes	yes	abstract
jorda_estimation_2005	§3	cite-only	yes	yes	abstract
jorda_local_nodate	§3	cite-only	yes	yes	abstract
dube	§3	cite-only	yes	yes	abstract
ramey_government_2018	§3	cite-only	yes	yes	abstract
[illegible]	[illegible]	[illegible]	yes	yes	[illegible]
[illegible]					

§1 Motivation

Vegh & Vuletin 2015 — How Is Tax Policy Conducted over the Business Cycle?

Citekey `@vegh_how_2015` · **Status** file ✓ · tag ✓ · §1 · consultation: abstract **Role** cite-only

Summary. Végh & Vuletin (2015, *AEJ: Economic Policy* 7(3):327–370) build a novel dataset of corporate-income, personal-income, and value-added *tax rates* for 62 countries over 1960–2013, explicitly to study the cyclicality of tax *policy* rather than tax revenues (which are endogenous to the cycle and therefore uninformative about policy intent). They find tax policy is acyclical in industrial countries but mostly procyclical in developing countries — mirroring the known pattern for government spending.

Relation. A foundational motivation cite: it establishes that statutory tax-rate cyclicality, not revenue cyclicality, is the policy-relevant object, and the authors themselves flag that the procyclical correlation "could just reflect the effect of tax multipliers" — precisely the exogeneity gap a [[#sec-romer2010]] narrative classification is built to resolve, and which our LLM-assisted pipeline aims to fill at scale across emerging markets. Links: [[#sec-romer2010]].

Jha et al. 2014 — Effectiveness of Countercyclical Fiscal Policy: Developing Asia

Citekey `@jha_effectiveness_2014` · **Status** file ✓ · tag ✓ · §1 · consultation: abstract **Role** cite-only

Summary. Jha, Mallick, Park & Quising (2014, *Journal of Macroeconomics* 40:82–98) examine whether discretionary fiscal policy stimulates output across 10 emerging Asian economies, identifying shocks via contemporaneous restrictions between fiscal and non-fiscal variables. Their most consistent finding is that in developing Asia tax cuts deliver a greater countercyclical output impact than government spending, implying scope for countercyclical tax adjustment where fiscal sustainability permits.

Relation. Beyond the tax-side effectiveness result, the paper underscores that most emerging-market revenue variation is non-discretionary (endogenous to the cycle) — exactly the endogeneity a narrative-identification approach removes, motivating why comparable exogenous tax-shock series are needed for credible EM fiscal analysis.

Frankel, Végh & Vuletin 2013 — On Graduation from Fiscal Procyclicality

Citekey `@frankel_graduation_2013` · **Status** file ✓ · tag ✓ · §1 · consultation: abstract **Role** cite-only

Summary. Frankel, Végh & Vuletin (2013, *Journal of Development Economics* 100(1):32–47) document that while industrial countries run countercyclical or acyclical fiscal policy and developing countries have historically been procyclical, about a third of the developing world "graduated" to countercyclical policy over the prior decade. Controlling for endogeneity, they identify a causal link from stronger institutional quality to less procyclical fiscal policy.

Relation. Supplies the "graduation" angle for the motivation section — raising whether the effectiveness and conduct of tax policy shifts as institutions improve, a question that comparable cross-country exogenous shock series would let researchers test directly. Filed in `_Fiscal Shocks` (Zotero `R64M3ZMY`).

Ilzetzki, Mendoza & Végh 2013 — How Big (Small?) Are Fiscal Multipliers?

Citekey `@ilzetzki_how_2013` · **Status** file ✓ · tag ✓ · §1 · consultation: abstract **Role** cite-only

Summary. Ilzetzki, Mendoza & Végh (2013, *Journal of Monetary Economics* 60(2):239–254) use a novel quarterly dataset of government expenditure in 44 countries to show that government-spending multipliers depend critically on country characteristics: the output effect is larger in industrial than developing countries, roughly zero under flexible exchange rates but large under predetermined regimes, smaller in open economies, and negative in high-debt countries.

Relation. The canonical spending-side multiplier-heterogeneity result; our Candidate-2 flagship is its tax-side analog, which requires comparable cross-country exogenous tax shocks of the kind our pipeline produces. Filed in `_Fiscal Shocks` (Zotero `XWMP6F75`).

Kaminsky, Reinhart & Végh 2004 — When It Rains, It Pours

Citekey `@kaminsky_when_2004` · Zotero `C375EN8P` · **Status** file ✓ · tag ✓ · §1 · consultation: abstract **Role** cite-only

Summary. Kaminsky, Reinhart & Végh (2004, *NBER Macroeconomics Annual* 19:11–53) document four stylized facts across 104 countries: net capital inflows are procyclical (especially for emerging markets), fiscal policy is procyclical in most developing countries, monetary policy appears procyclical in emerging markets, and capital-inflow periods coincide with expansionary macro policy — so for these countries, "when it rains, it pours."

Relation. The origin of the "when it rains, it pours" procyclicality puzzle our flagship decomposes; isolating exogenous fiscal shocks is what lets the procyclical correlation be interpreted causally rather than as a reflection of the cycle itself. Note: the former duplicate `QJWG8RSN` no longer appears in the collection (likely merged at source); citekey resolves in the re-exported bib.

Romer & Romer 2019 — Fiscal Space and the Aftermath of Financial Crises

Citekey `@romer_fiscal_2019` · Zotero `5JBXFRX9` · **Status** file ✓ · tag ✓ · §1 · consultation: abstract **Role** cite-only

Summary. Romer & Romer (2019, NBER Working Paper 25768) show that across OECD countries from 1980–2017, those with lower debt-to-GDP ratios responded to financial distress with far more expansionary fiscal policy and suffered milder aftermaths. They argue this fiscal-space relationship is driven partly by sovereign market access but even more by policymakers' choices, drawing on contemporaneous narrative accounts of the policy process.

Relation. Motivation-side framing for why fiscal space and the conduct of fiscal policy shape downturn severity; authored by the R&R team, but distinct from the tax-narrative identification method our pipeline automates.

Romer & Romer 2017 — New Evidence on the Aftermath of Financial Crises

Citekey `@romer_new_2017` · Zotero `HYNHB8X3` · **Status** file ✓ · tag ✓ · §1 · consultation: abstract **Role** cite-only

Summary. Romer & Romer (2017, *American Economic Review* 107(10):3072–3118) construct a new semiannual narrative series on financial distress for 24 OECD countries (1967–2012), classified on a fine scale from a consistent real-time source. They find output declines after financial crises are statistically significant and persistent but only moderate on average and highly variable across episodes, with the severity of distress itself a key driver of that variation.

Relation. Motivation-side framing companion on the output costs of financial crises in advanced economies; also a methodological exemplar of consistent real-time narrative classification, the lineage our automated approach extends. Links: [[#sec-rr_fiscal_space]].

Angrist & Pischke 2010 — The Credibility Revolution in Empirical Economics

Citekey `@angrist_credibility_2010` · **Status** file ✓ · tag ✓ · §1 · consultation: abstract **Role** cite-only

Summary. Angrist & Pischke (2010, *Journal of Economic Perspectives* 24(2):3–30) argue that empirical microeconomics has undergone a "credibility revolution" driven primarily by attention to research-design quality rather than sensitivity analysis. They note progress has been more dramatic in micro than in macroeconomics and industrial organization, and rebut the critique that design-driven work sacrifices good questions for clean answers.

Relation. The design-based framing for why credible identification matters, setting up the stakes for narrative identification — particularly the authors' own observation that macroeconomics lagged the credibility revolution, which our approach addresses.

Romer (Paul) — The Trouble With Macroeconomics

Citekey `@romer_trouble_nodate` · **Status** file ✓ · tag ✓ · §1 · consultation: abstract **Role** cite-only

Summary. Paul Romer's "The Trouble With Macroeconomics" (published version: 2016, *The American Economist*) polemicizes that for three decades macroeconomics has regressed: identification is no more credible than in the early 1970s but escapes challenge because it has become so opaque, and models attribute fluctuations to “imaginary” causal forces unmoved by any agent’s action. He likens the field’s deference to authority over fact to a general failure mode of science. (Zotero record has no date/venue populated.)

Relation. The sharp identification critique that a transparent narrative-identification-plus-validation approach is meant to answer — making the causal forces behind fiscal shocks explicit, fact-disciplined, and auditable rather than opaque.

Burke, Hsiang & Miguel 2015 — Global Non-linear Effect of Temperature on Economic Production

Citekey `@burke_global_2015` · Zotero `PU2P7CVS` · **Status** file ✓ · tag ✓ · §1 · consultation: abstract **Role** exemplar-other-domain

Summary. Burke, Hsiang & Miguel (2015, *Nature*) estimate a single, globally common non-linear relationship between annual temperature and aggregate economic production from a large subnational country panel, finding output peaks near 13 °C and falls sharply at higher temperatures — the same dose–response in rich and poor countries — then project large global GDP losses under continued warming.

Relation. The paper-architecture template: invent one new, comparable cross-country unit and estimate a single transferable function over it. That is the structural move our project makes by producing comparable *exogenous fiscal-shock* series across countries from heterogeneous national records. Imperfect §-fit (climate, not fiscal) — filed in §1 as a framing exemplar. Links: [[#sec-hsiang_quantifying_2013]].

Hsiang, Burke & Miguel 2013 — Quantifying the Influence of Climate on Human Conflict

Citekey `@hsiang_quantifying_2013` · Zotero `K4C3I9HW` · **Status** file ✓ · tag ✓ · §1 · consultation: abstract **Role** exemplar-other-domain

Summary. Hsiang, Burke & Miguel (2013, *Science*) meta-analyze dozens of studies across archaeology, criminology, economics, geography, history and political science onto a common standardized climate–conflict effect size, finding that warmer temperatures and rainfall extremes robustly raise both interpersonal and intergroup conflict.

Relation. The second Burke–Hsiang architecture template — converging heterogeneous studies onto one comparable unit — the rhetorical model for presenting a single cross-country fiscal-shock measure built from disparate national sources. Note: two Zotero copies exist (`K4C3I9HW`, `7C6BBH7P`); merge at source (citekey `hsiang_quantifying_2013` unaffected). Links: [[#sec-burke_global_2015]].

Acosta Ormaechea & Yoo 2012 — Tax Composition and Growth: A Broad Cross-Country Perspective

Citekey `@acosta_ormaechea_tax_2012` · **Status** file ✓ · tag ✓ · §1 · consultation: abstract **Role** cite-only

Summary. Acosta-Ormaechea & Yoo (2012, IMF Working Paper 12/257) provide a broad cross-country panel linking tax composition to long-run growth, finding that the mix of revenue instruments (income taxes versus consumption and property taxes) matters for growth outcomes. (No abstract in Zotero; summary grounded in title and venue — verify against source.)

Relation. The canonical "tax composition and growth" anchor the motivation needs, reinforcing that comparable cross-country tax-shock series are both scarce and valuable. Fills the §1 PENDING ACQUISITION slot in `lit_review.qmd`. Links: [[#sec-vegh_how_2015]].

Crispolti, Amaglobeli & Sheng 2022 — Cross-Country Evidence on the Revenue Impact of Tax Reforms

Citekey `@crispolti_cross-country_2022` · **Status** file ✓ · tag ✓ · §1 · consultation: abstract **Role** cite-only

Summary. Crispolti, Amaglobeli & Sheng (2022, IMF Working Paper 2022/199) estimate the revenue impact of discretionary tax-policy changes across 21 advanced and emerging economies using the IMF Tax Policy Reform Database v.4.0. Revenue yields vary by instrument and by the nature of the change — base-broadening of income, excise, and property taxes yields larger and longer-lasting revenue than rate changes, while rate changes matter more for VAT and social security contributions — and effects are asymmetric in direction.

Relation. Cross-country tax-reform/revenue evidence built on a structured tax-reform action dataset, motivating why comparable, instrument-level exogenous tax measures matter for emerging markets. **Dual-fit:** filed §1 (motivation) but defensibly §7 (adjacent to the IMF Tax Policy Reform Database / statutory-rate panels) — revisit at synthesis. Links: [[#sec-acosta_ormaechea_tax_2012]], [[#sec-dabla-norris_macroeconomic_2023]].

§2 Narrative identification (pillar)

Romer & Romer 2010 — Macroeconomic Effects of Tax Changes

Citekey `@romer2010` · **Status** file ✓ · tag ✓ · §2 · consultation: full **Role** read-deep

Summary. The foundational narrative method: read historical documents to identify *why* taxes changed and separate exogenous changes (long-run growth / deficit-reduction motivated) from endogenous ones (responding to current or prospective output). Full implementation detail lives in `docs/codebook_sources.md` §1.

Relation. The method we automate — the spine of the whole project. Every codebook (C1–C4) operationalizes a step of it. Links: [[#sec-romer]], [[#sec-cloyne_discretionary_2013]], [[#sec-romer_presidential_2023]].

Romer & Romer — A Narrative Analysis of Postwar Tax Changes

Citekey `@romer` · **Status** file ✓ · tag ✓ · §2 · consultation: full **Role** read-deep

Summary. The companion act-by-act narrative classification underlying `@romer2010`; the source of the 50-act reference index distilled in `docs/codebook_sources.md` §1.6.

Relation. Our gold reference for US validation (`us_shocks.csv` lineage). Links: [[#sec-romer2010]].

Romer & Romer 2023 — Presidential Address: Does Monetary Policy Matter?

Citekey `@romer_presidential_2023` · **Status** file ✓ · tag ✓ · §2 · consultation: abstract **Role** cite-only

Summary. In their AEA Presidential Address (Romer & Romer 2023, *American Economic Review* 113(6):1395–1423), the authors lay out the requirements for rigorous narrative macroeconomic identification, using monetary policy as the focal application. Reading the historical Minutes and Transcripts of Federal Reserve meetings (1946–2016), they identify policy changes *not* taken in response to current or prospective real-activity developments and find these monetary shocks have large, significant effects on unemployment, output, and inflation.

Relation. The canonical statement of the narrative-method doctrine, articulating the four requirements for credible narrative identification (qualitative source documents, careful causal reading, isolation of shocks orthogonal to current conditions, transparent classification). Das et al. adopt these four requirements (see `codebook_sources.md` §2.1), making this a doctrinal anchor for §2's narrative-identification framing even though the application here is monetary rather than fiscal.

Mertens & Ravn 2013 — Dynamic Effects of Personal and Corporate Income Tax Changes

Citekey `@mertens2013` · **Status** file ✓ · tag ✓ · §2 · consultation: abstract **Role** read-deep

Summary. Personal income tax changes are far more potent for output than corporate ones, a composition decomposition only the narrative record enables.

Relation. Motivates carrying PIT/CIT/VAT composition through our schema. Links: [[#sec-mertens_reconciliation_2014]].

Mertens & Ravn 2014 — A Reconciliation of SVAR and Narrative Tax-Multiplier Estimates

Citekey `@mertens_reconciliation_2014` · **Status** file ✓ · tag ✓ · §2 · consultation: abstract **Role** read-deep

Summary. Pure SVARs cannot recover the tax multiplier without assuming the output-elasticity of revenue; narrative proxies as instruments yield larger multipliers (~2 on impact, up to 3), and narrative measurement error is itself consequential.

Relation. Method-scrutiny anchor: why the narrative record is the necessary input, and a direct bridge to the §6 generated-regressor/measurement-error problem. Links: [[#sec-ludwig_large_2025]].

Cloyne 2013 — Discretionary Tax Changes and the Macroeconomy (UK)

Citekey `@cloyne_discretionary_2013` · **Status** file ✓ · tag ✓ · §2 · consultation: abstract **Role** read-deep

Summary. Applies the R&R narrative strategy to the UK and finds "remarkably similar" magnitudes (1% tax cut ⟶ +0.6% GDP on impact, +2.5% over three years).

Relation. The single most important precedent for our cross-country transfer claim — the only non-US tax-narrative series. Links: [[#sec-romer2010]].

Tillmann et al. 2024 — Disagreement about Fiscal Policy

Citekey `@tillmann_disagreement_2024` · **Status** file ✓ · tag ✓ · §2 · consultation: abstract **Role** cite-only

Summary. Tillmann, Winker, Latifi & Naboka-Krell (2024, SSRN working paper 5084219, "Disagreement about Fiscal Policy") apply NLP to the full set of German Bundestag speeches since 1960 to construct two fiscal-disagreement series from 1970 onward: government-vs-opposition and within-coalition disagreement. In a VAR, within-coalition disagreement has adverse macro effects — depressing industrial production, manufacturing orders, and stock prices — while government-opposition disagreement does not affect the business cycle.

Relation. Exemplifies the German textual-analysis tradition, which reads parliamentary *speech* rather than budget or treasury documents to measure fiscal conditions, contrasting with the document-based source corpus the R&R lineage relies on. Links: [[#sec-latifi_fiscal_nodate]].

Latifi et al. — Fiscal Policy in the Bundestag

Citekey `@latifi_fiscal_nodate` · **Status** file ✓ · tag ✓ · §2 · consultation: abstract **Role** cite-only

Summary. Latifi, Naboka-Krell, Tillmann & Winker ("Fiscal Policy in the Bundestag: Textual Analysis and Macroeconomic Effects," n.d.) use all Bundestag speeches (1960–2021) and an embedding-based approach that represents words and documents in a shared vector space to measure fiscal-policy sentiment on a contractionary-to-expansionary scale. Feeding fiscal sentiment into recursively-identified VARs, they show a change in sentiment causes a change in government spending with strong macroeconomic effects.

Relation. A contrast to the document-based R&R lineage we follow: rather than reading official budget/treasury narratives, it derives spending shocks from embedding-based sentiment over parliamentary debate. Links: [[#sec-tillmann_disagreement_2024]].

Kim 2015 — Understanding Narrative Inquiry

Citekey `@kim_understanding_2015` · **Status** file ✓ · tag ✓ · §2 · consultation: abstract **Role** cite-only

Summary. Kim (2015, *Understanding Narrative Inquiry: The Crafting and Analysis of Stories as Research*, SAGE) is a comprehensive introduction to narrative inquiry in the social and human sciences, covering the full process from situating narrative inquiry in its interdisciplinary and philosophical context to research design, data collection, analysis, interpretation, and theorizing narrative meaning.

Relation. Narrative-methodology background from the social-science (rather than econometric) tradition. A weak fit for §2's fiscal-shock focus, retained only as a framing cite for the broader meaning of "narrative" methodology.

Hebous & Zimmermann 2018 — Revisiting the Narrative Approach of Estimating Tax Multipliers

Citekey `@hebous_revisiting_2018` · **Status** file ✓ · tag ✓ · §2 · consultation: abstract **Role** cite-only

Summary. Hebous & Zimmermann (2018, *Scandinavian Journal of Economics* 120(2):428–439, "Revisiting the Narrative Approach of Estimating Tax Multipliers") test whether popular narrative tax-shock measures are relevant instruments for observable endogenous tax series. They find narrative tax measures are only weakly correlated with cyclically-adjusted tax revenues in the US and UK, and that under weak-instrument-robust inference the implied multipliers are often statistically insignificant — concluding the literature understates the uncertainty in narrative-approach multiplier estimates.

Relation. The "method scrutiny" cite that rounds out §2, probing the reliability of the R&R/narrative tax series and the multiplier magnitudes they imply, alongside [[#sec-mertens_reconciliation_2014]] — fills the §2 PENDING ACQUISITION slot in `lit_review.qmd`. Note: a duplicate Zotero record exists (`XGJJMI87` in-collection vs `VR6XHVG9`); merge at source. Links: [[#sec-romer2010]].

§3 Effects of fiscal policy

Blanchard & Perotti 2002 — Dynamic Effects of Spending and Taxes on Output

Citekey `@blanchard_empirical_nodate` · Zotero `XDDZP426` · **Status** file ✓ · tag ✓ · §3 · consultation: abstract **Role** cite-only

Summary. Blanchard & Perotti (2002, *Quarterly Journal of Economics* 117(4):1329–1368) develop a structural VAR / event-study approach that uses institutional information about tax and transfer programs and the timing of fiscal collections to identify the automatic response of fiscal variables to activity, isolating fiscal shocks. They find positive government-spending shocks raise output and positive tax shocks lower it, with spending shocks crowding out private investment. (No abstract in Zotero; summary from published QJE abstract.)

Relation. The SVAR baseline that identifies fiscal shocks via institutional-timing assumptions, and the methodological tradition our narrative approach offers an alternative to: rather than recovering shocks from a structural model's timing restrictions, we read the motivation for each measure directly from the documentary record. Citekey resolves in the re-exported bib.

Auerbach & Gorodnichenko 2012 — Measuring the Output Responses to Fiscal Policy

Citekey `@auerbach_measuring_2012` · Zotero `RKVTWKCU` · **Status** file ✓ · tag ✓ · §3 · consultation: abstract **Role** cite-only

Summary. Auerbach & Gorodnichenko (2012, *AEJ: Economic Policy* 4(2):1–27) use regime-switching models to estimate fiscal multipliers separately across recessions and expansions, finding spending multipliers considerably larger in recessions, that disaggregated spending (e.g., military) behaves differently from aggregate shocks, and that controlling for predictable components of fiscal shocks raises the recession multipliers.

Relation. Motivates state-dependence as a (de-prioritized) downstream question — multipliers that differ by cyclical regime — which our project treats as a payoff a credible exogenous shock series would feed into. Citekey resolves in the re-exported bib.

Gechert 2015 — What Fiscal Policy Is Most Effective? A Meta-Regression

Citekey `@gechert_what_2015` · Zotero `PN2C6PIY` · **Status** file ✓ · tag ✓ · §3 · consultation: abstract **Role** cite-only

Summary. Gechert (2015, *Oxford Economic Papers* 67(3):553–580, "What fiscal policy is most effective? A meta-regression analysis") synthesizes 104 multiplier studies, adjusting for study-design characteristics. Public-spending multipliers cluster near 1 and exceed tax/transfer multipliers by about 0.3–0.4 units, with public-investment multipliers larger still by roughly 0.5 unit; estimates vary systematically with study design.

Relation. A reference point for plausible multiplier magnitudes and for the spending-vs-tax composition gradient, anchoring expectations for what credible shock series should ultimately deliver downstream. Citekey resolves in the re-exported bib.

Riera-Crichton, Végh & Vuletin 2016 — Tax Multipliers: Pitfalls in Measurement

Citekey `@riera-crichton_tax_2016` · Zotero `TNSX9KDL` · **Status** file ✓ · tag ✓ · §3 · consultation: abstract **Role** cite-only

Summary. Riera-Crichton, Végh & Vuletin (2016, *Journal of Monetary Economics* 79:30–48, "Tax multipliers: Pitfalls in measurement and identification") build a novel 1980–2009 VAT-rate dataset and separate the identification problem (changes not triggered by output fluctuations) from the measurement problem (finding a policy variable under direct policymaker control). On identification their results favor narratives à la Romer & Romer (2010); on measurement they favor tax rates over revenue-based measures.

Relation. Catalogues the exact measurement and identification errors our magnitude/exogeneity outputs must avoid, and its endorsement of the R&R narrative for identification directly underwrites the method this project scales with LLMs. Citekey resolves in the re-exported bib.

Ramey 2019 — Ten Years after the Financial Crisis: Renaissance in Fiscal Research

Citekey `@ramey_ten_2019` · Zotero `SYE7PVL2` · **Status** file ✓ · tag ✓ · §3 · consultation: abstract **Role** cite-only

Summary. Ramey (2019, *Journal of Economic Perspectives* 33(2):89–114, "Ten Years after the Financial Crisis…") takes stock of post-crisis innovations in estimating fiscal multipliers, concluding most average-multiplier estimates lie in a narrow range (0.6–1 for spending, −2 to −3 for tax changes) while identifying circumstances that push multipliers outside those bands, and reviews the stimulus-vs-consolidation debate.

Relation. Surveys the post-crisis multiplier literature and thereby situates the demand for credible exogenous shock series — the input on which all of these estimates depend and which our pipeline aims to supply at scale. Citekey resolves in the re-exported bib.

Ramey 2016 — Macroeconomic Shocks and Their Propagation

Citekey `@ramey_macroeconomic_2016` · **Status** file ✓ · tag ✓ · §3 · consultation: abstract **Role** cite-only

Summary. Ramey (2016, *Handbook of Macroeconomics* vol. 2, "Macroeconomic Shocks and Their Propagation") reviews and synthesizes the identification of the shocks driving economic fluctuations, treating monetary, fiscal, and technology shocks in detail and presenting new estimates that reconcile competing parts of each literature.

Relation. The authoritative gap citation: the handbook's fiscal section is US-only, and the words "emerging"/"developing" never appear — precisely the coverage void our cross-country (US ⟶ Malaysia ⟶ SEA) extension targets.

Ramey 2011 — Identifying Government Spending Shocks: It's All in the Timing

Citekey `@ramey_identifying_2011` · **Status** file ✓ · tag ✓ · §3 · consultation: abstract **Role** cite-only

Summary. Ramey (2011, *Quarterly Journal of Economics* 126(1):1–50, "Identifying Government Spending Shocks: It's All in the Timing") shows that standard VAR-identified spending shocks are in fact anticipated by forward-looking agents, so the timing of identification matters decisively. Using a narrative "war dates" series capturing news about future spending, she finds responses (e.g., consumption and real wages falling) that differ from those obtained when anticipation is ignored. (No abstract in Zotero; summary from published QJE abstract.)

Relation. Establishes the anticipation/timing dimension that any shock series — and specifically our C3 (Timing) codebook — must respect, since misdating a measure relative to when it became known biases the recovered shock.

Jordà 2005 — Estimation and Inference of Impulse Responses by Local Projections

Citekey `@jorda_estimation_2005` · **Status** file ✓ · tag ✓ · §3 · consultation: abstract **Role** cite-only

Summary. Jordà (2005, *American Economic Review* 95(1):161–182, "Estimation and Inference of Impulse Responses by Local Projections") introduces local projections (LP): estimating impulse responses by direct horizon-by-horizon regression rather than iterating a fitted VAR. The method is computable with standard regression tools, more robust to misspecification, supports simple analytic inference, and accommodates highly nonlinear specifications.

Relation. The estimation toolkit our downstream payoff would use to map an identified shock series into dynamic output responses. Links: [[#sec-jorda_local_nodate]], [[#sec-dube]].

Jordà — Local Projections for Applied Economics

Citekey `@jorda_local_nodate` · **Status** file ✓ · tag ✓ · §3 · consultation: abstract **Role** cite-only

Summary. Jordà ("Local Projections for Applied Economics," n.d.) reviews how local projections connect the macro time-series language of VARs and impulse responses to the applied-micro language of potential outcomes and treatment effects, arguing LP serves as a natural bridge between the two traditions and identifying gains from greater methodological integration.

Relation. Bridges the macro IRF and applied-micro treatment-effects framings, clarifying that the impulse responses estimated from our shock series and micro causal-effect estimands are two dialects of the same estimator. Links: [[#sec-jorda_estimation_2005]].

Dube et al. — A Local Projections Approach to DiD Event Studies

Citekey `@dube` · **Status** file ✓ · tag ✓ · §3 · consultation: abstract **Role** cite-only

Summary. Dube, Girardi, Jordà & Taylor ("A Local Projections Approach to Difference-in-Differences Event Studies") propose an LP-DiD estimator that recasts dynamic heterogeneous-treatment-effect DiD problems using local projections plus a flexible "clean control" condition. The estimator is simple and transparent, avoids the negative-weighting problem of conventional fixed-effects DiD, admits arbitrary weighting and covariate controls, and performs well in simulations and two empirical applications.

Relation. A candidate downstream estimator for the shock series: it lets identified fiscal measures be treated as event-study "treatments" while sidestepping the bias problems of standard two-way fixed-effects designs.

Ramey & Zubairy 2018 — Government Spending Multipliers in Good Times and in Bad

Citekey `@ramey_government_2018` · **Status** file ✓ · tag ✓ · §3 · consultation: abstract **Role** cite-only

Summary. Ramey & Zubairy (2018, *Journal of Political Economy* 126(2):850–901, "Government Spending Multipliers in Good Times and in Bad…") use new quarterly historical US data spanning major wars and deep recessions to estimate spending multipliers by economic regime. They find multipliers below unity regardless of slack — robust across two identification schemes and two estimation methods — with more mixed, occasionally larger (up to ~1.5) estimates only at the zero lower bound.

Relation. Strengthens the state-dependence discussion in §3 by providing the contrarian finding (little evidence multipliers are larger in slack) and pairs the LP toolkit with a concrete fiscal application — fills the §3 PENDING ACQUISITION slot in `lit_review.qmd`. Note: two Zotero copies exist — the PDF copy `WRUHDRNP` (tagged, `duplicate-review`) and a no-PDF copy `46AXGD36`; merge at source keeping the PDF copy (citekey `ramey_government_2018` unaffected). Links: [[#sec-jorda_estimation_2005]].

Dabla-Norris & Lima 2023 — Macroeconomic Effects of Tax Rate and Base Changes

Citekey `@dabla-norris_macroeconomic_2023` · **Status** file ✓ · tag ✓ · §3 · consultation: abstract **Role** cite-only

Summary. Dabla-Norris & Lima (2023, *European Economic Review* 153:104399, "Macroeconomic effects of tax rate and base changes: Evidence from fiscal consolidations") build a new narrative dataset of 2,000+ tax measures across 10 OECD countries during consolidation years, coding expected yield, motivation, and announcement/implementation dates. Distinguishing tax-rate from tax-base changes (and across PIT, CIT, and VAT), they find base broadening causes smaller output and employment declines than rate hikes.

Relation. An effects-of-tax-changes study whose action-based narrative coding (yield, motivation, dates) closely parallels our pipeline's extraction targets, and whose rate-vs-base decomposition speaks directly to our magnitude/composition schema. **Dual-fit:** filed §3 (effects) but defensibly §7 (built on the action-based consolidation dataset lineage) — revisit at synthesis. Links: [[#sec-mertens2013]], [[#sec-crispolti_cross-country_2022]], [[#sec-pescatori_new_2011]].

§4 LLMs as measurement tools

Das et al. 2026 — Mapping Government Spending Actions Across 64 Countries

Citekey `@das_ai_2026` · Zotero `WBYNNC32` · **Status** file ✓ · tag ✓ · §4 · consultation: full **Role** competitor

Summary. Das, Furceri, Patel & Peralta-Alva (2026, IMF WP/26/43) build the first global *quarterly* narrative database of discretionary government-spending actions by applying a fixed, zero-shot GPT-4.1 prompt to Economist Intelligence Unit Country Reports (64 countries, 1952–2023; 8,636 episodes). The prompt screens R&R-style for *exogenous* spending actions, but the production output is a discrete direction proxy $z \in \{-1, 0, +1\}$ — spending-only, no tax side, no %GDP magnitude (an optional intensity field is excluded from the econometrics). Validation is fourfold: replicating R&R (2019) expert coding (>93%), a 51-run replicability test, local-projection prediction of realized spending, and ≥96% alignment with Adler et al. (2024) %GDP consolidations. Headline: a median cumulative spending multiplier ~0.7, larger in less-open economies, under fixed exchange rates, in downturns, and at the ZLB, identified via country-level Bayesian proxy-SVAR.

Relation. The flagship competitor — it shares our exact thesis that a fixed, country-agnostic, off-the-shelf LLM prompt can scale R&R narrative identification cross-country *without retraining*, and even shares design instincts we use (conservative net-aggregation of multi-action passages, verbatim-evidence retention, separate validation vs. production prompts). We differ on all four differentiators: (1) spending-only vs. our tax-side depth; (2) **direction only** — the authors explicitly "do not observe dollar-equivalent sizes … only their qualitative direction," and their 51-run test shows the granular intensity field is the *least* replicable while coarse direction is stable, corroborating our position that quantitative %GDP sizing is the hard part and our contribution; (3) no Halterman & Keith behavioral/zero-shot/error-analysis framework vs. our S0–S3; (4) English-only EIU reports vs. our EN/BM bilingual transfer. Their three-tier no-ground-truth validation (expert-replication % + aggregate prediction +

independent-dataset alignment) is a reusable template for our Malaysia/SEA setting. Note: Das positions only against Aruoba & Drechsel (monetary) and does *not* cite Bhasin or Fritsch — the competitor cluster is one we draw, not one they assert. Links: [[#sec-bhasin_identifying_2026]], [[#sec-fritsch_high-frequency_nodate]], [[#sec-adler_2024]], [[#sec-aruoba_identifying_nodate]].

Bhasin & Loungani 2026 — Identifying Episodes of Fiscal Austerity

Citekey `@bhasin_identifying_2026` · Zotero `BLQ3YEEH` · **Status** file ✓ · tag ✓ · §4 · consultation: full **Role** competitor

Summary. Bhasin & Loungani (2026, *J. Economic Dynamics & Control*) build a hierarchical five-prompt LLM cascade to identify fiscal-austerity (consolidation) episodes from IMF Article IV staff reports (2004–2020, 17 OECD countries): country filter ⟶ fiscal-policy gate ⟶ four-pillar austerity definition ⟶ R&R-style exogeneity filter ⟶ %GDP magnitude. Core model Llama-3.3-70B (GPT-4.1 / DeepSeek-R1 robustness checks converge). Validated against the Guajardo et al. (2014) / Adler et al. (2024) consolidation dataset over 289 country-years: precision 0.93, recall 0.90, F1 0.91. Headline: LLM-identified shocks yield *smaller* multipliers than the narrative benchmark, and a swap-decomposition attributes the divergence to shock *sizing*, not event selection.

Relation. The closest direct competitor on method — LLM-automated R&R identification with an explicit exogeneity gate *and* a %GDP magnitude step — sharing our cross-country ambition. It diverges on our differentiators: consolidation episodes broadly (revenue *or* expenditure lumped into a binary austerity label) vs. our tax decomposition; magnitude secondary to a binary label (size correlation only 0.59, systematically smaller than benchmark) vs. our first-class per-act %GDP; a single harmonized *secondary* source (Article IV, written *about* countries by IMF staff), validated against existing ground truth, vs. our *primary* budget/treasury documents with no Malaysia/SEA ground truth and H&K/expert-agreement gates they do not use. **Correction to the prior seed:** the paper does say "shock sizing, not event selection," but frames it *oppositely* — it argues the *narrative benchmark's* larger multipliers are inflated by persistence/residual endogeneity (ρ≈0.73, predictable from lagged GDP) and presents its own smaller, less-persistent sizes as *better-identified.* So it concedes only that Article IV reports yield conservative ("headline-package") magnitudes, not that magnitude is a binding accuracy constraint — our positioning should not overstate this. Their Prompt-4 exogeneity exemplars (Maastricht/SGP compliance, constitutional debt-brake, pre-announced medium-term programs) are concrete worked examples that could sharpen our C2 exogenous-side examples. Links: [[#sec-das_ai_2026]], [[#sec-fritsch_high-frequency_nodate]], [[#sec-guajardo_leigh_pescatori]], [[#sec-adler_2024]].

Fritsch & Mazlish — High-Frequency Fiscal Shocks

Citekey `@fritsch_high-frequency_nodate` · Zotero `FWILZBHC` · **Status** file ✓ · tag ✓ · §4 · consultation: full **Role** competitor

Summary. Fritsch & Mazlish ("High-Frequency Fiscal Shocks," Oxford, preliminary) build a two-agent LLM pipeline producing daily-frequency US fiscal shocks (1947–2025): a *forecaster* agent reads each day's fiscal/macro news (full Reuters/ WSJ articles for 2005–25; summaries pre-2005) and updates a present-value-of-deficits forecast, while a *classifier* agent labels each change endogenous / fiscal-event / both, apportions it into purchases (G), taxes/transfers (T) and interest (I), and tags permanence and stabilization. It thus recovers signed, %GDP-scaled magnitudes *with* a tax-vs-spending split at daily/event frequency, with a prompt-imposed knowledge-cutoff guarding against look-ahead bias. Validation: correlation 0.71 with Ramey (2011) defense shocks and **0.89 with R&R (2010) tax shocks**, plus high-frequency asset-price reactions (10Y yields, dollar, term spread) and multipliers.

Relation. The most direct competitor on three of our differentiators at once — signed %GDP magnitude, R&R-style exogeneity, and a tax-side decomposition — and its 0.89 correlation with R&R tax shocks is a strong validation benchmark. Our defensible distinctions: (i) it is **US-only**, with cross-country deployment merely *promised* as future work, whereas our contribution is the *executed* country-agnostic transfer (US⟶Malaysia⟶SEA) with frozen codebooks and expert-agreement gates; (ii) its inputs are **news wires** building a *forecast/ expectations* series, fundamentally different from our reading of *primary* budget/ treasury documents to classify *enacted* measures; (iii) it lacks the formal H&K 5-stage scaffolding we use, relying on correlation-recovery and asset-price/ multiplier validation. Its high-frequency asset-price identification is *orthogonal* to our document-based approach (our unit is the enacted act at quarterly/annual frequency, not a datestamped market event), so it is a complement to cite, not a method to adopt. Links: [[#sec-das_ai_2026]], [[#sec-bhasin_identifying_2026]].

Aruoba & Drechsel — Identifying Monetary Policy Shocks: A Natural-Language Approach

Citekey `@aruoba_identifying_nodate` · Zotero `NP8WKD26` · **Status** file ✓ · tag ✓ · §4 · consultation: abstract **Role** exemplar-other-domain

Summary. Aruoba & Drechsel identify monetary-policy shocks by applying natural-language-processing techniques to the documents Federal Reserve staff prepare in advance of policy decisions, extracting the information set behind each decision and purging the systematic component; the residual surprises are validated by the absence of forecast-error predictability and by theory-consistent impulse responses.

Relation. The validate-without-ground-truth precedent we follow: measurement validity is established through forecast-error orthogonality and theory-consistent IRFs rather than concordance against an existing series — the template for our no-ground-truth Phase 2 gates. Das et al. (2026) likewise position against this as their main narrative-identification precedent. Note: two Zotero copies (`NP8WKD26` journalArticle, no date; `I43BYWTD` 2024 preprint); merge at source. Links: [[#sec-asirvatham_gpt_2026]], [[#sec-das_ai_2026]].

Fernandez-Fuertes 2025 — Monetary Policy Shocks: LLMs and Central-Bank Communication

Citekey `@fernandez-fuertes_monetary_2025` · Zotero `QBBU9ZQ2` · **Status** file ✓ · tag ✓ · §4 · consultation: abstract **Role** exemplar-other-domain

Summary. Fernandez-Fuertes (2025, SSRN preprint) develops a multi-agent LLM framework that reads Fed communications (Beige Books, Minutes) released before each FOMC meeting to generate conditional expectations and construct narrative monetary-policy surprises that are *less noisy* than market-based measures, yielding theoretically consistent impulse responses and profitable yield-curve strategies. By extracting expectations directly rather than cleaning surprises ex post, it implements narrative identification at scale without high-frequency contamination.

Relation. A multi-agent-LLM-as-measurement deployment outside fiscal policy, and the closest monetary-side analog to our pipeline: it shares the move of extracting measurements/expectations directly from policy documents, and its explicit contamination-avoidance concern echoes the threat our S1 memorization test guards against. A companion to Aruoba & Drechsel on the monetary side of the LLM-narrative-identification frontier. Links: [[#sec-aruoba_identifying_nodate]], [[#sec-asirvatham_gpt_2026]].

Torreblanca et al. 2025 — The Credibility Revolution in Political Science

Citekey `@torreblanca_credibility_2025` · Zotero `NUBCCS2W` · **Status** file ✓ · tag ✓ · §4 · consultation: abstract **Role** exemplar-other-domain

Summary. Torreblanca, Dinneen, Grossman & Xu (2025, preprint) use an LLM to classify 91,632 political-science articles published 2003–2023 across 174 journals by causal research design, charting how the credibility revolution reshaped the discipline.

Relation. A large-scale LLM-as-classifier exemplar in a social science — demonstrating reliable LLM annotation at the corpus scale our deployment operates in, and pairing the LLM measurement tool with the credibility-revolution framing (cf. Angrist & Pischke) that motivates our identification stakes. Links: [[#sec-angrist_credibility_2010]].

Gentzkow, Kelly & Taddy 2019 — Text as Data

Citekey `@gentzkow_text_2019` · **Status** file ✓ · tag ✓ · §4 · consultation: abstract **Role** cite-only

Summary. Gentzkow, Kelly & Taddy (2019, *Journal of Economic Literature* 57(3):535–574, “Text as Data”) provide the canonical introduction to using text as an input to economic research: what makes text distinct from other data, a practical statistical overview centered on the high-dimensionality problem (the vocabulary far exceeds the sample), and a survey of applications across economics.

Relation. The canon of text-as-data for economics, and its central lesson — that text measures are model-dependent and that fitted text representations notoriously fail to transfer across domains and corpora — is precisely the adversarial constraint our frozen, country-agnostic codebooks must survive. It frames the measurement-validity problem that [[#sec-halterman2025]] operationalizes into a staged validation discipline. Links: [[#sec-halterman2025]].

Korinek 2025 — AI Agents for Economic Research

Citekey `@korinek_ai_2025` · **Status** file ✓ · tag ✓ · §4 · consultation: abstract **Role** cite-only

Summary. Korinek (2025, NBER Working Paper 34202, "AI Agents for Economic Research") is a hands-on guide to building autonomous LLM-based agents that plan, use tools, and execute multi-step research tasks — literature reviews, writing and debugging econometric code, fetching and analyzing data, coordinating workflows — aimed at making agentic frameworks accessible to economists, and traces AI's evolution from chatbots to reasoning models to agents.

Relation. The LLM-for-economists meta layer for §4: it legitimizes building LLM agents into the research workflow while cataloguing the practical caveats — hallucination, error cascades across multi-step chains, and prompt brittleness — that any deployed pipeline must guard against, stopping short of vouching for LLMs in causal inference itself. Pairs with [[#sec-dell_deep_2025]].

Feuerriegel et al. 2025 — Using NLP to Analyse Text Data in Behavioural Science

Citekey `@feuerriegel_using_2025` · **Status** file ✓ · tag ✓ · §4 · consultation: abstract **Role** cite-only

Summary. Feuerriegel et al. (2025, *Nature Reviews Psychology* 4(2):96–111, "Using natural language processing to analyse text data in behavioural science") review the full NLP pipeline for behavioural-science text analysis, walking through modelling choices from dictionary-based approaches to large language models, foregrounding the interpretability-versus-accuracy trade-off and closing with recommendations for rigour and reproducibility.

Relation. Situates our tooling within the broader NLP-for-social-science methods literature, supplying the pipeline-level vocabulary (modelling choices, the interpretability/accuracy trade-off, reproducibility discipline) that our codebook-driven LLM annotation instantiates for the fiscal-policy domain.

Ash & Hansen 2023 — Text Algorithms in Economics

Citekey `@ash_text_2023` · **Status** file ✓ · tag ✓ · §4 · consultation: abstract **Role** cite-only

Summary. Ash & Hansen (2023, *Annual Review of Economics* 15:659–688, "Text Algorithms in Economics") survey algorithmic text analysis in economics around three contributions: methods for representing documents and words as high-dimensional and embedding vectors, four core empirical tasks that encompass most text-as-data research, and a candid flagging of current limitations — with particular emphasis on the challenge of validating algorithmic output.

Relation. The econ text-as-data canon alongside [[#sec-gentzkow_text_2019]], situating our LLM tooling within the methods lineage; its explicit worry about validating algorithmic text output is exactly the gap an H&K-style validation discipline is meant to close. Fills part of the §4 PENDING ACQUISITION slot.

Grimmer, Roberts & Stewart 2022 — Text as Data: A New Framework for Machine Learning and the Social Sciences

Citekey `@grimmer_text_2022` · **Status** file ✓ (book; no PDF) · tag ✓ · §4 · consultation: abstract **Role** cite-only

Summary. Grimmer, Roberts & Stewart (2022, Princeton University Press, *Text as Data: A New Framework for Machine Learning and the Social Sciences*) is the textbook framework organizing text-as-data research around core tasks — representation, discovery, measurement, prediction, and causal inference — via a sequential, iterative, inductive research design, stressing that measurement from text demands explicit validation rather than assumed validity. (Book; no PDF.)

Relation. The measurement-validity framing our H&K pipeline instantiates — the discovery ⟶ measurement ⟶ inference workflow and the validation discipline measurement requires — a canonical §4 cite. The validation standard it sets is what [[#sec-halterman2025]] turns into concrete behavioral and zero-shot tests for LLM coders. Citekey resolves in `references.bib`. Links: [[#sec-halterman2025]].

Gilardi, Alizadeh & Kubli 2023 — ChatGPT Outperforms Crowd Workers for Text-Annotation Tasks

Citekey `@gilardi_chatgpt_2023` · **Status** file ✓ · tag ✓ · §4 · consultation: abstract **Role** cite-only

Summary. Gilardi, Alizadeh & Kubli (2023, *PNAS* 120(30):e2305016120, "ChatGPT outperforms crowd workers for text-annotation tasks") show that on four samples of tweets and news articles (n=6,183), zero-shot ChatGPT beats MTurk crowd workers on relevance, stance, topic, and frame detection by about 25 percentage points on average, with higher intercoder agreement than both crowd workers and trained annotators, at roughly one-thirtieth the per-annotation cost.

Relation. The LLM-annotation-validity anchor for §4 — the social-science precedent that LLMs can stand in for human coders — which our expert-agreement gates both rely on and stress-test in a higher-stakes domain. It motivates the very validation discipline that [[#sec-halterman2025]] formalizes before such substitution can be trusted. Links: [[#sec-halterman2025]].

Ziems et al. 2024 — Can Large Language Models Transform Computational Social Science?

Citekey `@ziems_can_2024` · **Status** file ✓ · tag ✓ · §4 · consultation: abstract
Role cite-only

Summary. Ziems et al. (2024, *Computational Linguistics* 50(1):237–291, "Can Large Language Models Transform Computational Social Science?") benchmark 13 LLMs zero-shot across 25 representative English CSS tasks, contributing prompting best practices and an evaluation pipeline. They find LLMs do not beat the best fine-tuned classifiers on taxonomic labeling but reach fair human agreement, while their free-form explanations often exceed crowdworker gold references — concluding LLMs can augment, not replace, the CSS pipeline.

Relation. The capabilities-and-limits survey that bounds the LLM-as-annotator claim our pipeline rests on, mapping where zero-shot models can and cannot substitute for human coders. It tempers the strong-form results of [[#sec-gilardi_chatgpt_2023]] and [[#sec-tornberg_large_2025]] by documenting the classification tasks where fine-tuned models still win. Links: [[#sec-gilardi_chatgpt_2023]], [[#sec-tornberg_large_2025]].

Törnberg 2025 — LLMs Outperform Expert Coders and Supervised Classifiers at Annotating Political Social Media

Citekey `@tornberg_large_2025` · **Status** file ✓ · tag ✓ · §4 · consultation: abstract
Role cite-only

Summary. Törnberg (2025, *Social Science Computer Review* 43(6):1181–1195, "Large Language Models Outperform Expert Coders and Supervised Classifiers at Annotating Political Social Media Messages") finds zero-shot GPT-4 beats both supervised classifiers and human coders (expert and crowd) at identifying politicians' political affiliation from single tweets across 11 countries — 0.934 accuracy and 0.982 inter-coder reliability in the US — succeeding even on cases requiring interpretation of implicit references and contextual reasoning.

Relation. A strong-form LLM-annotation-validity cite reinforcing [[#sec-gilardi_chatgpt_2023]] — the precedent our expert-agreement gates lean on — and its multilingual, multi-country result is directly relevant to our cross-country transfer ambition. Its strong claims are bounded by the more cautious task-by-task findings of [[#sec-ziems_can_2024]]. Links: [[#sec-ziems_can_2024]].

Dell 2025 — Deep Learning for Economists

Citekey `@dell_deep_2025` · **Status** file ✓ · tag ✓ · §4 · consultation: abstract
Role cite-only

Summary. Dell (2025, *Journal of Economic Literature* 63(1):5–58, "Deep Learning for Economists") reviews deep neural methods for imputing structured information from large-scale unstructured text and image corpora — classifiers, regression and generative models, and embeddings — with applications spanning classification, document digitization, record linkage, and data exploration, arguing such models are cheap to tune and scale affordably to millions or billions of observations.

Relation. The LLM-for-economists meta layer situating our tooling within the discipline's methods canon, complementing the agent-building emphasis of [[#sec-korinek_ai_2025]] and the foundational text-as-data framing of [[#sec-gentzkow_text_2019]] with a structured-extraction-and-scaling lens that matches our act-level identification task. Links: [[#sec-korinek_ai_2025]], [[#sec-gentzkow_text_2019]].

Gambacorta et al. 2024 — CB-LMs: Language Models for Central Banking

Citekey `@gambacorta_cb-lms_2024` · **Status** file ✓ · tag ✓ · §4 · consultation: abstract **Role** cite-only

Summary. Gambacorta, Kwon, Park, Patelli & Zhu (2024, BIS Working Papers No. 1215, "CB-LMs: language models for central banking") introduce central-bank language models — specialised encoder-only models retrained on a corpus of central-bank speeches, policy documents, and research papers — and show they outperform their foundational base models at predicting masked words in central-bank idioms.

Relation. A domain-specialized-LM exemplar for applying language models to government/policy economic text, demonstrating measurable gains from adapting models to institutional fiscal/monetary idiom. Pairs with the institutional overview in [[#sec-araujo_artificial_2024]]. Links: [[#sec-araujo_artificial_2024]].

Araujo et al. 2024 — Artificial Intelligence in Central Banking

Citekey `@araujo_artificial_2024` · **Status** file ✓ · tag ✓ · §4 · consultation: abstract **Role** cite-only

Summary. Araujo, Doerr, Gambacorta & Tissot (2024, BIS Bulletin No. 84, "Artificial intelligence in central banking") survey AI and LLM adoption across central-bank functions, including the analysis of policy and economic documents alongside data management, supervision, and macro/financial monitoring. (No abstract in Zotero; summary grounded in the BIS Bulletin metadata and known scope — verify against source.)

Relation. The BIS LLM-primer the plan names for §4, situating LLM-based economic-text analysis institutionally within central-bank practice. Pairs with [[#sec-gambacorta_cb-lms_2024]], the domain-specialized-model counterpart from the same BIS group. Links: [[#sec-gambacorta_cb-lms_2024]].

§5 Validation (pillar)

Halterman & Keith 2025 — Codebook LLMs

Citekey `@halterman2025` · **Status** file ✓ · tag ✓ · §5 · consultation: full **Role** read-deep

Summary. The five-stage codebook-LLM measurement framework (prepare ⟶ behavioral tests ⟶ zero-shot eval ⟶ error analysis ⟶ fine-tune), built on the premise that codebook conceptualization is a first-order concern. Distilled in `docs/codebook_sources.md` §3.

Relation. The validation spine of the project — our S0–S3 pipeline is its instantiation. Links: [[#sec-asirvatham_gpt_2026]], [[#sec-gentzkow_text_2019]].

Asirvatham, Mokski & Shleifer 2026 — GPT as a Measurement Tool

Citekey `@asirvatham_gpt_2026` · **Status** file ✓ · tag ✓ · §5 · consultation: full **Role** read-deep

Summary. The GABRIEL package; separates **accuracy** (benchmark agreement) from **directness** (the measurement comes from the text, via a signal-stripping test); contamination-via-cutoff and prompt-robustness evidence; human-parity framing. Distilled in `docs/codebook_sources.md` §5.2.

Relation. Closest external template for "is the LLM output a valid measurement"; seeds a proposed C2 motivation-masking robustness test and the human-parity framing for our no-ground-truth Phase 2 gates. Links: [[#sec-halterman2025]], [[#sec-ludwig_large_2025]].

Rathje et al. 2024 — GPT for Multilingual Psychological Text Analysis

Citekey `@rathje_gpt_2024` · Zotero `8LABL4VM` · **Status** file ✓ · tag ✓ · §5 · consultation: abstract **Role** exemplar-other-domain

Summary. Rathje et al. (2024, *PNAS Nexus*) show that GPT detects psychological constructs (sentiment, discrete emotions, offensiveness) across 12 languages with accessible zero-shot prompts, outperforming established dictionary methods and approaching fine-tuned models without per-language retraining.

Relation. Directly relevant to our EN/BM cross-language stress test: evidence that a single LLM can transfer a measurement construct across languages with no per-language retraining — the capability our Malaysia bilingual leg depends on — while the cross-lingual benchmark literature warns that this transfer degrades for lower-resource languages. Links: [[#sec-ahuja_mega_2023]], [[#sec-singh_indicgenbench_2024]], [[#sec-xuan_mmlu-prox_2025]], [[#sec-halterman2025]].

Singh 2026 — Deadly Stigma

Citekey `@singh_deadly_2026` · Zotero `C3AVEYPB` · **Status** file ✓ · tag ✓ · §5 · consultation: abstract **Role** exemplar-other-domain

Summary. Singh (2026, preprint "Deadly Stigma") uses a physician-validated LLM to detect stigmatizing language in clinical notes, then exploits exogenous variation in stigma to estimate its real-world harm — solving the measurement problem (stigma is hard to observe) so that a credible research design becomes feasible.

Relation. An LLM-as-measurement validated against expert (physician) judgment and then fed into a causal design — the same two-step pattern as our pipeline (LLM measurement ⟶ downstream inference), with expert validation as the credibility anchor. A close analog to our expert-agreement gates and to the §6 generated-regressor concern of carrying LLM measurement error into downstream estimates. Links: [[#sec-zhao_agentic_2026]], [[#sec-halterman2025]].

Zhao et al. 2026 — An Agentic System for Rare-Disease Diagnosis

Citekey `@zhao_agentic_2026` · Zotero `ATXC3JZM` · **Status** file ✓ · tag ✓ · §5 · consultation: abstract **Role** exemplar-other-domain

Summary. Zhao et al. (2026, *Nature*) present a multi-agent LLM system for rare-disease diagnosis with *traceable reasoning* chains, reporting ~95% expert agreement on its reasoning steps (not just final labels) and strong diagnostic accuracy across a large patient cohort.

Relation. The traceable-reasoning + expert-agreement validation pattern: experts verify the reasoning that produced a label, not only the label itself. A direct model for how our expert-validation gates could assess C2 motivation *rationales* (the evidence-and-reasoning behind an exogenous/endogenous call), strengthening the credibility argument beyond marginal-label agreement. Links: [[#sec-singh_deadly_2026]].

Fellegi & Sunter 1969 — A Theory for Record Linkage

Citekey `@fellegi_theory_1969` · **Status** file ✓ · tag ✓ · §5 · consultation: abstract **Role** cite-only

Summary. Fellegi & Sunter (1969, *JASA* 64(328):1183–1210) develop the foundational mathematical model for record linkage: deciding whether a comparison-pair from two files represents the same entity. Each pair is assigned link / possible link / non-link based on a comparison vector γ of coded agreements, with the positive dispositions governed by stipulated error probabilities over the matched/unmatched likelihoods m(γ) and u(γ). Their central theorem constructs an *optimal* linkage rule that minimizes the probability of failing to make a positive disposition for fixed error levels.

Relation. The classical anchor for the C0 aggregation / entity-resolution subtask — the "is this the same act" problem our canonical clustering solves with an LLM rather than a likelihood-ratio rule, though the matched/non-matched decision under controlled error rates is the same target. Fills the §5-core record-linkage PENDING ACQUISITION slot. Links: [[#sec-abramitzky_automated_2021]], [[#sec-binette_almost_2022]].

Abramitzky, Boustan & Eriksson et al. 2021 — Automated Linking of Historical Data

Citekey `@abramitzky_automated_2021` · **Status** file ✓ · tag ✓ · §5 · consultation: abstract **Role** cite-only

Summary. Abramitzky, Boustan, Eriksson, Feigenbaum & Pérez (2021, *Journal of Economic Literature* 59(3):865–918, "Automated Linking of Historical Data") evaluate automated methods for linking individuals across digitized census records against each other and against hand-linking. They report that several automated methods achieve very low (<5%) false-positive rates along a frontier trading off false positives against match rate, that humans and algorithms agree closely on the same linking variables, and that downstream coefficient estimates are stable across linkers.

Relation. The econ-history precedent for automated linking that legitimizes a machine-assisted C0 aggregation step — its false-positive/match-rate frontier and human-versus-algorithm agreement framing directly parallel how we trade off over-merging against fragmentation. It operationalizes the Fellegi–Sunter foundation for economics. Links: [[#sec-fellegi_theory_1969]], [[#sec-binette_almost_2022]].

Binette & Steorts 2022 — (Almost) All of Entity Resolution

Citekey `@binette_almost_2022` · **Status** file ✓ · tag ✓ · §5 · consultation: abstract **Role** cite-only

Summary. Binette & Steorts (2022, *Science Advances* 8(12):eabi8021, "(Almost) All of Entity Resolution") survey structured entity resolution (record linkage / deduplication), motivated by applications from census population estimation and conflict-death counts to author disambiguation. They review seminal and modern probabilistic/Bayesian methods drawn from statistics, computer science, ML, databases, economics, and political science, and close with practical research directions.

Relation. The state-of-the-art ER framing for C0: its blocking → pairwise-matching → clustering → evaluation-against-ground-truth vocabulary maps directly onto our year-blocked and canonical-clustering methods, situating the classical decision-theoretic foundation and the econ-history linking benchmarks within a unified modern taxonomy. Links: [[#sec-fellegi_theory_1969]], [[#sec-abramitzky_automated_2021]].

Carlini et al. 2022 — Quantifying Memorization Across Neural Language Models

Citekey `@carlini_quantifying_2022` · **Status** file ✓ · tag ✓ · §5 · consultation: full **Role** read-deep

Summary. Measures verbatim memorization by prompting models with true training-set prefixes and checking exact-suffix reproduction under greedy decoding. Establishes three log-linear laws: extractable memorization rises with model size (≈+19 pp per 10× parameters; 2–5× across a family), with training-data duplication, and with prompt-context length — the "discoverability" effect, 33%⟶65% extractable from 50⟶450 tokens of context. A GPT-2-vs-GPT-Neo control (prompt a model on data it could *not* have seen and compare) isolates genuine memorization from generic predictive competence, and estimates ≥1% of The Pile is verbatim-extractable from a 6B model.

Relation. The mechanistic backing for our H&K **S1 memorization test**. The threat is concrete: the published 44-act R&R list and the source ERP/Treasury documents are exactly the high-value, plausibly-duplicated public text most at risk of verbatim recall, so a model could "recognize" an act by name/year rather than reason from the supplied chunk. Carlini supplies both the risk levers (larger models, duplication, long name-and-year-rich prefixes) and the control logic our S1 probe borrows — withhold or perturb the identifying surface forms (act name, P.L. number, year) and check whether the classification survives. Because memorization scales with model size, S1 must be re-run when the deployment model changes, not treated as a one-time check. Pillar deep-read distilled into `docs/codebook_sources.md` §3. Links: [[#sec-halterman2025]], [[#sec-golchin_time_2024]], [[#sec-deng_unveiling_2024]].

Golchin & Surdeanu 2024 — Time Travel in LLMs: Tracing Data Contamination

Citekey `@golchin_time_2024` · **Status** file ✓ · tag ✓ · §5 · consultation: abstract
Role cite-only

Summary. Golchin & Surdeanu (2024, *ICLR 2024*, "Time Travel in LLMs: Tracing Data Contamination in Large Language Models") propose prompt-based methods to detect whether specific datasets leaked into an LLM's pretraining corpus, probing for instance- and partition-level contamination by prompting the model to reproduce or complete dataset content (guided by metadata cues such as split names) and comparing generations against the reference. (No abstract in Zotero; summary from the verified title/venue and the established method.)

Relation. A contamination-detection technique relevant to certifying that our US historical corpus is *read, not recalled* — a concern because R&R-era US fiscal documents are plausibly in pretraining data, which would inflate apparent skill. Pairs with the memorization-quantification line and the broader contamination survey. Links: [[#sec-carlini_quantifying_2022]], [[#sec-deng_unveiling_2024]].

Deng et al. 2024 — Unveiling the Spectrum of Data Contamination: A Survey

Citekey `@deng_unveiling_2024` · **Status** file ✓ · tag ✓ · §5 · consultation: abstract
Role cite-only

Summary. Deng et al. (2024, *Findings of ACL 2024*, "Unveiling the Spectrum of Data Contamination in Language Model: A Survey from Detection to Remediation") present the first comprehensive survey of data contamination — the overlap between LLM training corpora and evaluation benchmarks — examining its effects across stages and forms, categorizing detection methods by focus, assumptions, strengths, and limitations, and reviewing mitigation strategies.

Relation. The survey anchor for the contamination/memorization risk our H&K S1 test guards against, providing the taxonomy by which to situate our certification that the model analyzes rather than recalls the US corpus, and framing the specific detection techniques we draw on. Links: [[#sec-carlini_quantifying_2022]], [[#sec-golchin_time_2024]].

Ahuja et al. 2023 — MEGA: Multilingual Evaluation of Generative AI

Citekey `@ahuja_mega_2023` · **Status** file ✓ · tag ✓ · §5 · consultation: abstract **Role** cite-only

Summary. Ahuja et al. (2023, *EMNLP 2023*, “MEGA: Multilingual Evaluation of Generative AI”) present the first comprehensive benchmarking of generative LLMs (including ChatGPT and GPT-4) across 16 NLP datasets spanning 70 typologically diverse languages, comparing them to SOTA non-autoregressive models and documenting systematic performance shortfalls for generative models on low-resource languages relative to English.

Relation. Evidence of the cross-lingual performance gap our EN/BM (and SEA-extension) stress test must measure — the documented English-versus-low-resource degradation is exactly the failure mode we probe for in Bahasa Malaysia. Anchors the cross-lingual-risk evidence base. Links: [[#sec-singh_indicgenbench_2024]], [[#sec-xuan_mmlu-prox_2025]], [[#sec-rathje_gpt_2024]].

Singh et al. 2024 — IndicGenBench: Multilingual Benchmark for LLM Generation

Citekey `@singh_indicgenbench_2024` · **Status** file ✓ · tag ✓ · §5 · consultation: abstract **Role** cite-only

Summary. Singh et al. (2024, *ACL 2024*, “IndicGenBench”) release the largest benchmark for evaluating LLM generation across 29 Indic languages (13 scripts, 4 families), covering cross-lingual summarization, machine translation, and cross-lingual QA, with human-curated multi-way parallel data. Evaluating GPT-3.5, GPT-4, PaLM-2, and LLaMA, they find a significant performance gap in all languages relative to English even for the strongest models.

Relation. A low-resource-language evaluation precedent directly analogous to the Bahasa Malaysia leg of our cross-language test, and its multi-way parallel construction is a model for assessing like-for-like degradation. Reinforces the cross-lingual gap evidence base. Links: [[#sec-ahuja_mega_2023]], [[#sec-xuan_mmlu-prox_2025]].

Xuan et al. 2025 — MMLU-ProX: A Multilingual Benchmark for Advanced LLM Evaluation

Citekey `@xuan_mmlu-prox_2025` · **Status** file ✓ · tag ✓ · §5 · consultation: abstract **Role** cite-only

Summary. Xuan et al. (2025, *EMNLP 2025*, "MMLU-ProX: A Multilingual Benchmark for Advanced Large Language Model Evaluation") introduce a 29-language benchmark of 11,829 *identical* questions per language (plus a 658-question lite version), built via multi-LLM translation with expert review expressly to enable direct cross-lingual comparison. Evaluating 36 SOTA LLMs, they find strong performance in high-resource languages but marked declines in low-resource ones, especially African languages.

Relation. Its parallel-question design (identical items per language) is the methodological template for a like-for-like EN/BM comparison, isolating language effect from item difficulty; reinforces the cross-lingual-risk evidence base. Links: [[#sec-ahuja_mega_2023]], [[#sec-singh_indicgenbench_2024]].

§6 Generated-regressor inference (pillar)

Ludwig, Mullainathan & Rambachan 2025 — LLMs: An Applied Econometric Framework

Citekey `@ludwig_large_2025` · **Status** file ✓ · tag ✓ · §6 · consultation: full **Role** read-deep

Summary. Splits LLM uses into prediction (needs "no training leakage") and estimation/measurement (needs a validation sample for valid downstream inference). Accuracy ≠ validity (bias = correlation of LLM error with regressors); debias against a small random validation sample (125–250 items help); model/prompt choice can flip downstream estimates absent debiasing. Distilled in `docs/codebook_sources.md` §5.1.

Relation. A §6 pillar anchor — our LLM fiscal shocks are exactly their "estimation problem" feeding a VAR/LP. The de-risking trio (Battaglia, Angelopoulos, Egami) is now filed, ending the single-source risk. Links: [[#sec-mertens_reconciliation_2014]], [[#sec-asirvatham_gpt_2026]], [[#sec-battaglia_inference_2024]], [[#sec-angelopoulos_prediction-powered_2023]], [[#sec-egami_using_2023]].

Battaglia, Christensen, Hansen & Sacher 2024 — Inference for Regression with Variables Generated from Unstructured Data

Citekey `@battaglia_inference_2024` · **Status** file ✓ · tag ✓ · §6 · consultation: full **Role** read-deep

Summary. Formalizes inference for the two-step strategy in which an upstream model estimates latent variables from text (e.g. topic shares) and a downstream regression treats them as observed. Proves this is a measurement-error problem whose bias scales with the average inverse information per observation ($\kappa = \lim \sqrt{n} \cdot E[1/C_i]$): when $\kappa > 0$ the OLS estimator is consistent but its asymptotic distribution is *re-centered*, so ordinary (Eicker–Huber–White) standard errors have the right *width* but the wrong *center* and under-cover; only $\kappa = 0$ gives valid two-step inference. Their baseline correction is a "one-step" joint upstream+downstream likelihood estimated by HMC — which presumes a generative IR likelihood an LLM judge does not provide. Extends explicitly to VARs/IRFs and similarity-measure regressions.

Relation. The most directly on-point external source for §6: our LLM-measured shocks are exactly a "variable generated from unstructured data" feeding a VAR/LP, and their VAR/IRF extension is our downstream object. Two transfer caveats: their κ formula keys the bias to document length under a topic-model likelihood (our LLM-classification error is not literally a $1/C_i$ sampling process — the qualitative lesson transfers, the exact formula does not), and their joint-MLE fix needs a likelihood the LLM lacks (the authors flag this as the open frontier). Because we *have* a validation sample (44 US labeled acts; expert-labeled Malaysia), the operative correction for us is the labeled-subset strand they contrast against, with Battaglia supplying the bias *diagnosis*: the error mode is mis-centering, not mis-sizing, so it cannot be detected by inspecting reported standard errors. Pillar deep-read distilled into `docs/codebook_sources.md` §5.4. Links: [[#sec-ludwig_large_2025]], [[#sec-egami_using_2023]], [[#sec-angelopoulos_prediction-powered_2023]], [[#sec-mertens_reconciliation_2014]].

Angelopoulos, Bates, Fannjiang, Jordan & Zrnic 2023 — Prediction-Powered Inference

Citekey `@angelopoulos_prediction-powered_2023` · **Status** file ✓ · tag ✓ · §6 · consultation: full **Role** read-deep

Summary. Prediction-powered inference (PPI) builds provably valid confidence intervals and P-values for a population estimand from a small gold-labeled sample (n) plus a large unlabeled sample ($N \gg n$) carrying ML predictions. It combines a *measure of fit* on the large imputed set with a *rectifier* — the prediction error estimated only on the gold sample — to debias the imputation; coverage holds for *any* ML algorithm and *any* data distribution, with no assumption on model quality. Intervals are strictly narrower than gold-only when predictions are decent and N is large; the framework covers means, quantiles, OLS/logistic coefficients (any convex-objective estimand), plus covariate- and label-shift variants. It offers no gain when predictions are too weak or N/n too small.

Relation. The general, assumption-light machinery underwriting our Phase 2 design: a small expert-labeled subset (PPI's gold sample) plus a large LLM-labeled corpus (the prediction sample) is a textbook PPI instance. It converts the expert subset from a one-off agreement check into the rectifier that yields valid downstream confidence intervals over the LLM-measured shock series — even when the LLM is biased, which is exactly our cross-country, no-retraining situation. The naive "treat LLM labels as truth" path is the imputation estimator the paper shows fails to cover; expert-only is valid but underpowered; PPI is the principled middle. Its stated limitation (no gain when predictions are weak or N not large relative to n) is a live risk at SEA corpus sizes, and a `ppi-py` reference implementation exists. Links: [[#sec-egami_using_2023]], [[#sec-battaglia_inference_2024]], [[#sec-ludwig_large_2025]], [[#sec-asirvatham_gpt_2026]].

Egami, Hinck, Stewart & Wei 2023 — Using Imperfect Surrogates for Downstream Inference

Citekey `@egami_using_2023` · **Status** file ✓ · tag ✓ · §6 · consultation: full **Role** read-deep

Summary. Design-based supervised learning (DSL) uses imperfect, possibly-biased LLM/ML "surrogate" labels as outcomes in downstream regressions while preserving valid inference. Surrogates are observed for all n documents, but expert gold labels only for a *probability-sampled* subset drawn with a known sampling probability π bounded away from zero; DSL forms a doubly-robust pseudo-outcome (cross-fitted prediction + inverse-propensity residual correction on the gold rows) and substitutes it into the downstream moment equations. Because π is known *by design*, estimates are consistent with valid confidence intervals *regardless* of how biased the surrogate is — surrogate quality affects only efficiency. Naive surrogate-only use gave ~40% coverage of nominal-95% CIs at 90% accuracy in their simulations; DSL matched gold-only bias/coverage at materially lower RMSE across 18 datasets.

Relation. The social-science-facing formalization closest to our expert-validation workflow: LLM labels over the full act corpus, expert labels on a validated subset, feeding downstream macro inference. Its load-bearing requirement reshapes our protocol — the validated subset must be *probability-sampled* with a recorded π (simple or stratified), not a convenience or borderline-acts sample — and in return the expert validation does double duty: not just agreement measurement (our ≥80%/≥70% targets) but active debiasing, with SEs that survive arbitrarily biased LLM labels. This reframes the agreement target as an *efficiency* lever rather than a validity gate. Distinct from PPI in exploiting researcher-controlled sampling (removing the nuisance-rate requirement), which our fully-controlled corpus naturally satisfies. Scope caveat: DSL targets the known-corpus outcome-regression case, so our US⟶Malaysia cross-country transfer sits at the edge of its domain-shift scope. Links: [[#sec-angelopoulos_prediction-powered_2023]], [[#sec-battaglia_inference_2024]], [[#sec-ludwig_large_2025]].

§7 Datasets

Devries, Guajardo, Leigh & Pescatori 2011 — A New Action-Based Dataset of Fiscal Consolidation

Citekey `@pescatori_new_2011` · Zotero `C9YM26FG` · **Status** file ✓ · tag ✓ · §7 · consultation: abstract **Role** benchmark

Summary. Devries, Guajardo, Leigh & Pescatori (2011, IMF WP) assemble the original action-based narrative dataset of fiscal-consolidation measures for 17 OECD economies, reading contemporaneous policy documents to identify changes motivated by deficit reduction rather than by a response to prospective economic conditions.

Relation. The founding narrative-comparator dataset that Guajardo et al. (2014) and Adler et al. (2024) extend, and the lineage Das and Bhasin validate against — the benchmark class our output landscape is measured against. Fills the §7 PENDING ACQUISITION slot in `lit_review.qmd`. Links: [[#sec-guajardo_leigh_pescatori]], [[#sec-adler_2024]].

Adler et al. 2024 — An Updated Action-Based Dataset of Fiscal Consolidation

Citekey `@adler_updated_2024` · Zotero `7ZQA6BJ3` · **Status** file ✓ · tag ✓ · §7 · consultation: abstract **Role** benchmark

Summary. Adler, Allen, Ganelli & Leigh (2024, IMF WP 2024/210) extend the action-based fiscal-consolidation dataset to a longer, broader OECD panel, supplying %GDP-sized consolidation measures identified from contemporaneous policy documents.

Relation. The OECD %GDP benchmark that the LLM competitors lean on — Das et al. report ≥96% alignment against it and Bhasin validates against the same lineage — so it is the quantitative-magnitude comparator our per-act %GDP output is positioned beside. Links: [[#sec-guajardo_leigh_pescatori]], [[#sec-pescatori_new_2011]], [[#sec-das_ai_2026]], [[#sec-bhasin_identifying_2026]].

Guajardo, Leigh & Pescatori 2014 — Expansionary Austerity? International Evidence

Citekey `@guajardo_expansionary_2014` · Zotero `MFGJ6A8H` · **Status** file ✓ · tag ✓ · §7 · consultation: abstract **Role** benchmark

Summary. Guajardo, Leigh & Pescatori (2014, *JEEA*) use the action-based narrative dataset to estimate the short-run effects of fiscal consolidation on OECD output, finding consolidation is contractionary for private demand and GDP — and that conventional cyclically-adjusted measures bias estimates toward the "expansionary austerity" view.

Relation. The narrative-comparator study that demonstrates *why* exogenous (action-based) identification overturns conclusions drawn from endogenous fiscal measures — the exact bias our exogeneity classification removes — and the benchmark Bhasin validates against. Links: [[#sec-adler_2024]], [[#sec-pescatori_new_2011]], [[#sec-bhasin_identifying_2026]].

Romer & Romer 2019 — Replication Data for the 2010 Tax-Shock Measure

Citekey `@romer_replication_2019` · Zotero `YCM7BMYI` · **Status** file ✓ · tag ✓ · §7 · consultation: abstract **Role** benchmark

Summary. Romer & Romer (2019) ICPSR/openICPSR replication archive for the 2010 tax-shock measure — the act-by-act narrative classification and constructed exogenous tax-change series underlying "The Macroeconomic Effects of Tax Changes."

Relation. The US reference series (`us_shocks.csv` lineage) our pipeline replicates and validates against end-to-end in Phase 1, and the source the §4 LLM competitors (Das, Fritsch) benchmark their recovery against. Links: [[#sec-romer2010]].

References