

Literature review

Esteban Degetau
World Bank
estebandegetau@gmail.com

Agustín Samano
World Bank
asamanopenaloza@worldbank.org

June 24, 2026

Table of contents

1. Motivation: why exogenous fiscal measures matter	2
2. Pillar: the narrative identification approach we are scaling	2
3. How these shocks get used: effects of fiscal policy	3
4. LLMs as measurement tools: text-as-data foundations and deployment	4
5. Pillar: is the LLM output good enough to be a measurement?	5
6. Pillar: using LLM-generated variables in downstream inference	6
7. Existing fiscal datasets: the output landscape and comparators	7
Open sequencing note	7
References	8

Role: the synthesized, full-length review prose for the methodological paper, organized by the argument in `docs/lit_review_plan.md` (§1–§7). It draws on the tagged `_Fiscal Shocks` Zotero library; the positioning argument it serves lives in `docs/brainstorm.qmd`, and the codebook-implementation distillations in `docs/codebook_sources.md`. As of the latest BetterBibTeX re-export every cited source now resolves to a `@-citekey` in `references.bib`; no `PENDING ACQUISITION` markers or plain-text “pending re-export” placeholders remain.

The methodological spine is four moves: take a hand-built narrative identification method (§2), scale it with LLMs (§4), validate the LLM as a measurement tool (§5), and hand off a generated variable for downstream inference (§6). Sections 2, 5, and 6 are the load-bearing pillars; sections 1, 3, and 7 motivate and situate.

Table 1: Coverage of the seven review sections in the current library.

Section	Role	Items	Status
§1 Motivation	Situate	13	Complete
§2 Narrative	Pillar	10	Strong
§3 Effects	Situate	12	Complete
§4 LLM text-as-data	Situate	17	Competitors + canon complete
§5 Validation	Pillar	14	Core gaps closed
§6 Generated-regressor	Pillar	4	De-risking trio filed
§7 Datasets	Situate	4	Complete

1. Motivation: why exogenous fiscal measures matter

The “why.” For most emerging markets we lack consistent, comparable measures of exogenous fiscal shocks, the input credible fiscal analysis needs. Kept compressed; it does not lead the paper.

The motivating puzzle is fiscal procyclicality in the developing world: poor countries’ fiscal policies tend to amplify rather than smooth their own crises. The stylized fact originates with Kaminsky et al. (2004) on “when it rains, it pours,” and Vegh and Vuletin (2015) sharpen it on the tax side, documenting that tax *policy*, not just tax revenue, is procyclical in developing countries while acyclical in industrial ones. Crucially, they concede that the procyclical correlation “could just reflect the effect of tax multipliers” because their statutory-rate series carries no record of *motivation*, and they name the Christina D. Romer and Romer (2010) narrative classification as the missing input. This is the gap our pipeline fills.

The development-macro stakes are well established. Ilzetzki et al. (2013) show fiscal multipliers depend on country characteristics (exchange-rate regime, openness, debt), Frankel et al. (2013) document that institutional quality governs whether a country “graduates” from procyclicality, and Jha et al. (2014) find that in developing Asia tax cuts have larger countercyclical effects than spending, while conceding that roughly 70% of revenue variation is non-discretionary, the exact endogeneity that narrative motivation removes. The same team’s work on fiscal space and the aftermath of financial crises (Christina D. Romer and Romer 2019a, 2017) underscores that the conduct of fiscal policy materially shapes downturn severity, raising the stakes for measuring it credibly. The broader identification stakes are framed by the credibility revolution (Angrist and Pischke 2010) and, in macro specifically, by the critique that macroeconomic identification remains opaque (P. Romer, n.d.).

What makes comparable cross-country shock series both scarce and valuable is the tax-composition–growth channel: Acosta Ormaechea et al. (2012) anchor the long-run link between the revenue mix and growth, and Crispolti et al. (2022) show, on a structured tax-reform database, that revenue yields vary systematically by instrument and by whether a reform changes rates or bases — exactly the instrument-level detail an exogenous-shock series must carry. Methodologically, the paper’s architecture follows the Burke–Hsiang template of inventing one comparable cross-country unit and estimating a single transferable function over it (Burke et al. 2015; Hsiang et al. 2013).

2. Pillar: the narrative identification approach we are scaling

The method we automate. The methodological backbone, distinct from “effects” (§3) and “datasets” (§7).

The foundation is Christina D. Romer and Romer (2010) and its companion narrative analysis (Christina D. Romer and Romer, n.d.), which read historical

documents to identify *why* taxes changed and to separate exogenous tax changes (motivated by long-run growth or deficit-reduction goals) from endogenous ones (responding to current or prospective output). The methodological doctrine, what makes narrative analysis rigorous, is stated most explicitly in Christina D. Romer and Romer (2023). These are distilled for codebook design in [docs/codebook_sources.md](#) §1.

The method travels. Cloyne (2013) applies the R&R strategy to the United Kingdom and finds “remarkably similar” magnitudes (a 1% tax cut raises GDP about 0.6% on impact and 2.5% over three years), the single most important precedent for our cross-country transfer claim. On the composition margin, Mertens and Ravn (2013) show personal income tax changes are far more potent for output than corporate ones, a decomposition only the narrative record enables.

The method is also contested, which is why §2 must engage scrutiny rather than only foundation. Mertens and Ravn (2014) prove that pure SVARs cannot recover the tax multiplier without assuming the output-elasticity of revenue, that narrative proxies used as instruments deliver larger multipliers (around 2 on impact, up to 3), and that measurement error in narrative series is itself consequential, a point that links directly to the generated-regressor problem in §6. Hebous and Zimmermann (2018) push the scrutiny further: the popular narrative tax measures are only weakly correlated with cyclically-adjusted revenues in the US and UK, and under weak-instrument-robust inference the implied multipliers are often insignificant, so the literature understates narrative-approach uncertainty. A parallel textual-analysis tradition in Germany (Latifi et al., n.d.; Tillmann et al. 2024) reads parliamentary speech rather than budget documents, a useful contrast to the document-based R&R lineage we follow. Narrative-inquiry methodology more broadly is surveyed by Kim (2015).

3. How these shocks get used: effects of fiscal policy

Shows the demand for the variable we produce, and motivates the §6 pillar (what happens when the regressor is LLM-measured).

The demand side is the multiplier literature. The SVAR tradition begins with Blanchard and Perotti (n.d.) and the narrative-vs-SVAR reconciliation runs through Mertens and Ravn (2014). State-dependence, multipliers that differ in recessions and expansions, is established by Auerbach and Gorodnichenko (2012), while Ramey and Zubairy (2018) supply the contrarian finding that spending multipliers stay below unity regardless of slack — so the regime-dependence is itself contested. The measurement pitfalls specific to tax multipliers are catalogued by Riera-Crichton et al. (2016), whose endorsement of R&R-style narrative identification directly underwrites the method we scale. The magnitude of multipliers across studies is synthesized by the meta-regression of Gechert (2015) and by Ramey’s surveys (Ramey 2016, 2019). Ramey (2011) adds the anticipation/timing dimension that any fiscal-shock series must respect. On the tax side specifically, Dabla-Norris and Lima (2023) build a narrative dataset of 2,000+ consolidation-era tax measures and find base-broadening less contrac-

tionary than rate hikes — its yield/motivation/date coding closely parallels our own extraction targets.

The estimation toolkit our downstream payoff would use is local projections (Jordà 2005, n.d.), including the difference-in-differences event-study variant (Dube et al., n.d.). The authoritative gap citation closing this section is Ramey (2016): the fiscal portion of the handbook is essentially US-only, and the words “emerging” and “developing” do not appear.

4. LLMs as measurement tools: text-as-data foundations and deployment

Situates the tooling in the econ text-as-data canon and the social-science annotation literature, and introduces the direct competitors.

The econ text-as-data canon is anchored by Gentzkow et al. (2019), who also supply an adversarial flag we must answer: fixed text measures notoriously fail to transfer across domains, exactly the property our frozen, country-agnostic codebooks claim. Ash and Hansen (2023) survey the same lineage and foreground the challenge of *validating* algorithmic text output, the gap an H&K-style discipline closes, and Grimmer et al. (2022) give the book-length statement that measurement from text demands explicit validation rather than assumed validity. The meta layer for economists runs from Korinek (2025), who catalogues hallucination, error cascades, and prompt brittleness and pointedly declines to vouch for LLMs in causal inference, to the structured-extraction-and-scaling lens of Dell (2025); the NLP pipeline itself is reviewed for social scientists by Feuerriegel et al. (2025).

The social-science annotation literature supplies the precedent that LLMs can substitute for human coders. Gilardi et al. (2023) show zero-shot ChatGPT beats crowd workers on text annotation by about 25 points with higher intercoder agreement, and Törnberg (2025) that it beats both expert coders and supervised classifiers at multi-country political-affiliation labeling. Ziems et al. (2024) bound the claim, documenting taxonomic-labeling tasks where fine-tuned models still win and concluding LLMs augment rather than replace the pipeline. This is the validity our expert-agreement gates both lean on and stress-test in a higher-stakes, primary-document domain.

The direct competition is three LLM-based narrative fiscal papers that, read together, share one revealing shortcut. Das et al. (2026) map spending actions across 64 countries from EIU Country Reports but recover only qualitative direction, no %GDP magnitude, and defer taxes. Bhasin and Loungani (2026) extract austerity episodes from IMF Article IV reports for 17 OECD countries; they do attribute their LLM-versus-benchmark divergence to “shock sizing, not event selection,” but frame it *oppositely* to a magnitude-failure reading — arguing their smaller, less-persistent sizes are better-identified than an inflated narrative benchmark — so our positioning should treat magnitude as the hard, contested step, not claim they concede a binding accuracy limit. Fritsch and Mazlish (n.d.) recover daily %GDP fiscal surprises with a tax-vs-spending split

(0.89 correlation with R&R tax shocks), but for the US only and from news wires building an *expectations* series. Each reads a secondary, harmonized, English-language source precisely to avoid the cost of primary national documents, and each validates by concordance against a pre-existing narrative series, a strategy that cannot travel to the countries the method is meant to serve. The (non-LLM) monetary analog Aruoba and Drechsel (n.d.) points to the alternative: validate against forecast-error predictability and theory-consistent impulse responses rather than an existing label set. Related LLM-as-measurement deployments outside fiscal policy include central-bank-communication analysis (Fernandez-Fuertes 2025) and the classification of research designs across 91,632 political-science articles (Torreblanca et al. 2025); the institutional uptake of LLMs across central-bank functions is surveyed by Araujo et al. (2024) and instantiated in domain-specialized models by Gambacorta et al. (2024).

5. Pillar: is the LLM output good enough to be a measurement?

The H&K spine plus the broader measurement-validity tradition, organized by our three sub-tasks: extraction, aggregation/entity-resolution, categorization.

The framework is Halterman and Keith (2025): a five-stage pipeline (prepare a codebook for humans and LLMs, test basic capabilities, evaluate zero-shot accuracy, analyze errors, and only then fine-tune) built on the premise that “codebook conceptualization is still a first-order concern.” It is distilled for our behavioral tests and error analysis in [docs/codebook_sources.md](#) §3, and it supplies the central insight that off-the-shelf LLMs follow real codebooks imperfectly zero-shot.

The closest external validation template is Asirvatham et al. (2026) (the GABRIEL package), which separates two requirements our project conflates at its peril: **accuracy** (agreement with a benchmark) and **directness** (the measurement comes from the text, not from a correlated proxy). Their signal-stripping test, remove the rated content and confirm the rating collapses (a 98% reduction in their case), is the model for a high-value robustness check we should add to C2: mask the stated motivation and re-classify exogenous-versus-endogenous; if the call survives, the model is inferring from act name, era, or party rather than reading the rationale. They also reframe “is the model good enough” as human-parity (model-versus-consensus compared to leave-one-out human-versus-consensus), the correct frame for our no-ground-truth Phase 2 expert-agreement gates, and endorse the held-out discipline our S2/S3 split already enforces. Exemplars of LLM-as-measurement validated against expert judgment in other domains reinforce the pattern: multilingual psychological annotation (Rathje et al. 2024), which is directly relevant to our EN/BM stress test; physician-validated detection of stigmatizing clinical language (M. Singh 2026); and a traceable-reasoning diagnostic agent with 95.4% expert agreement (Zhao et al. 2026).

Two further sub-tasks of our pipeline draw on literatures beyond H&K. **Aggregation / entity resolution (C0)**, recognizing that “the Goods and Services Tax”

and “Cukai Barang dan Perkhidmatan” are one act, is the classical record-linkage problem: anchored decision-theoretically by Fellegi and Sunter (1969), given its modern taxonomy by Binette and Steorts (2022), and shown to be reliably automatable in economics by Abramitzky et al. (2021), whose false-positive/match-rate frontier and human-versus-algorithm agreement directly parallel the over-merging-versus-fragmentation trade-off our canonical clustering manages. **Contamination / memorization**, the threat that historical US documents and the published R&R act list sit in pretraining and let a model recognize an act by name rather than reason from the supplied chunk, is quantified mechanistically by Carlini et al. (2022): extractable memorization rises with model size, training-data duplication, and prompt-context length, so the S1 memorization test must be re-run whenever the deployment model changes rather than treated as one-off. Detection techniques come from Golchin and Surdeanu (2024) and the surveying taxonomy from Deng et al. (2024). The **cross-lingual** risk our EN/BM and SEA legs face is documented by the multilingual-evaluation literature: Ahuja et al. (2023), H. Singh et al. (2024), and Xuan et al. (2025) all find marked degradation on lower-resource languages relative to English, and the latter two supply the parallel-item design (identical questions per language) that isolates the language effect from item difficulty in a like-for-like EN/BM comparison.

6. Pillar: using LLM-generated variables in downstream inference

The riskiest gap. Our output becomes a regressor in a VAR or local projection, making it a generated-regressor / measurement-error problem.

This is the pillar with the least coverage and the highest referee risk. The anchor is Ludwig et al. (2025), who frame precisely our case: an LLM-measured variable feeding a downstream causal estimate is an “estimation problem,” and valid inference requires combining LLM outputs with a small labeled validation sample, not merely high classification accuracy. Three of their results bind on us, and are distilled in [docs/codebook_sources.md](#) §5:

- **Accuracy is not validity.** A model uniformly close to the truth can still produce arbitrarily biased downstream estimates; the bias equals the correlation of the LLM’s error with the regressors. For us the dangerous regressor is the business cycle, the exact axis exogeneity hinges on.
- **Debias against a validation sample.** Do not plug raw LLM shocks into the multiplier regression; use the debiased estimator against a small random expert-labeled sample, which can yield tighter standard errors than the expert sample alone. Their simulations suggest 125–250 validation items suffice, a number infeasible at 20–40 Malaysian acts, forcing an honest design choice (fraction-of-corpus or pooled-regional validation, or a scope-limited inferential claim).
- **Specification search is a real risk.** Absent debiasing, the choice of model or prompt can change downstream coefficients in sign and significance, which directly implicates our Haiku-to-Sonnet C0 decision and any codebook iteration.

Three formal treatments of inference with a generated regressor now end this pillar’s single-source exposure. Battaglia et al. (2024) is the most directly on point: they prove the two-step “estimate a latent variable from text, then regress on it” strategy is a measurement-error problem whose bias *re-centers* the estimator’s asymptotic distribution, so ordinary (Eicker–Huber–White) standard errors have the right width but the wrong center and under-cover — and their explicit VAR/IRF extension is exactly our downstream object. The operative lesson is diagnostic: because the error mode is mis-centering, not mis-sizing, it cannot be caught by inspecting reported standard errors. Angelopoulos et al. (2023) (prediction-powered inference) and Egami et al. (2023) (design-based supervised learning) both convert a small gold-labeled subset plus a large LLM-labeled corpus into valid downstream inference *regardless* of model quality — PPI via a rectifier estimated on the gold sample, DSL via a doubly-robust pseudo-outcome over a *probability-sampled* gold subset. DSL’s requirement that the validated subset be probability-sampled with a recorded sampling rate (not a convenience or borderline-acts sample) directly reshapes our expert-validation protocol and reframes the $\geq 80\%$ / $\geq 70\%$ agreement targets as efficiency levers, not merely validity gates. Their shared limitation — little gain when predictions are weak or the unlabeled corpus is not large relative to the gold sample — is a live risk at 20–40 Malaysian acts, and the honest design constraint Phase 2 must confront. These are distilled for our protocol in [docs/codebook_sources.md](#) §5.

7. Existing fiscal datasets: the output landscape and comparators

What already exists, so our contribution and validation targets are legible.

The narrative / action-based lineage is the relevant comparator class. Its origin is the action-based fiscal-consolidation dataset of Pescatori et al. (2011), extended by Guajardo et al. (2014) (which shows action-based identification overturns the “expansionary austerity” reading that conventional cyclically-adjusted measures produce) and updated with %GDP sizing by Adler et al. (2024) — the OECD benchmarks that Das et al. (2026) and Bhasin and Loungani (2026) validate against. On the statutory / text-mined side, Crispolti et al. (2022) build on the IMF Tax Policy Reform Database and Vegh and Vuletin (2015) supply the statutory-rate panel. The US reference series we replicate and validate against is the Romer-Romer shock list, available as replication data (Christina D. Romer and Romer 2019b) and underlying [us_shocks.csv](#).

Open sequencing note

§1 (motivation) and §3 (effects) are confirmed as motivation and demand framing, not application targets. Pillars §2, §5, and §6 carry the methodological contribution; with the generated-regressor trio (Battaglia, Angelopoulos, Egami) now filed and the §5 record-linkage, contamination, and cross-lingual cites resolved, all three pillars are multiply sourced and no **PENDING ACQUISITION** marker remains.

References

- Abramitzky, Ran, Leah Boustan, Katherine Eriksson, James Feigenbaum, and Santiago Pérez. 2021. “Automated Linking of Historical Data.” *Journal of Economic Literature* 59 (3): 865–918.
- Acosta Ormaechea, Santiago, SAcosta Ormaechea@imf.org, Jiae Yoo, and JYoo@imf.org. 2012. “Tax Composition and Growth: A Broad Cross-Country Perspective.” *IMF Working Papers* 12 (257): 1.
- Adler, Gustavo, Cian Allen, Giovanni Ganelli, and Daniel Leigh. 2024. “An Updated Action-Based Dataset of Fiscal Consolidation.” *IMF Working Papers* 2024 (210): 1.
- Ahuja, Kabir, Harshita Diddee, Rishav Hada, et al. 2023. “MEGA: Multilingual Evaluation of Generative AI.” In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, edited by Houda Bouamor, Juan Pino, and Kalika Bali, 4232–67. Singapore: Association for Computational Linguistics.
- Angelopoulos, Anastasios N., Stephen Bates, Clara Fannjiang, Michael I. Jordan, and Tijana Zrnica. 2023. “Prediction-Powered Inference.” *Science* 382 (6671): 669–74.
- Angrist, Joshua D., and Jörn-Steffen Pischke. 2010. “The Credibility Revolution in Empirical Economics: How Better Research Design Is Taking the Con Out of Econometrics.” *Journal of Economic Perspectives* 24 (2): 3–30.
- Araujo, Douglas, Sebastian Doerr, Leonardo Gambacorta, and Bruno Tissot. 2024. *Artificial Intelligence in Central Banking*. nos. No 84 (January).
- Aruoba, S. Boragan, and Thomas Drechsel. n.d. “Identifying Monetary Policy Shocks: A Natural Language Approach.” *American Economic Journal: Macroeconomics*.
- Ash, Elliott, and Stephen Hansen. 2023. “Text Algorithms in Economics.” *Annual Review of Economics* 15 (Volume 15, 2023): 659–88.
- Asirvatham, Hemanth, Elliott Moksiki, and Andrei Shleifer. 2026. *GPT as a Measurement Tool*. Working {Paper}. Working Paper Series. National Bureau of Economic Research.
- Auerbach, Alan J., and Yuriy Gorodnichenko. 2012. “Measuring the Output Responses to Fiscal Policy.” *American Economic Journal: Economic Policy* 4 (2): 1–27.
- Battaglia, Laura, Timothy Christensen, Stephen Hansen, and Szymon Sacher. 2024. *Inference for Regression with Variables Generated from Unstructured Data*. {SSRN} {Scholarly} {Paper}. Rochester, NY: Social Science Research Network.
- Bhasin, Karan, and Prakash Loungani. 2026. “Identifying Episodes of Fiscal Austerity: An LLM-Based Approach.” *Journal of Economic Dynamics and Control* 188 (July): 105351.

- Binette, Olivier, and Rebecca C. Steorts. 2022. “(Almost) All of Entity Resolution.” *Science Advances* 8 (12): eabi8021.
- Blanchard, Olivier, and Roberto Perotti. n.d. *AN EMPIRICAL CHARACTERIZATION OF THE DYNAMIC EFFECTS OF CHANGES IN GOVERNMENT SPENDING AND TAXES ON OUTPUT*.
- Burke, Marshall, Solomon M. Hsiang, and Edward Miguel. 2015. “Global Non-Linear Effect of Temperature on Economic Production.” *Nature* 527 (7577): 235–39.
- Carlini, Nicholas, Daphne Ippolito, Matthew Jagielski, Katherine Lee, Florian Tramèr, and Chiyuan Zhang. 2022. “Quantifying Memorization Across Neural Language Models.” September.
- Cloyne, James. 2013. “Discretionary Tax Changes and the Macroeconomy: New Narrative Evidence from the United Kingdom.” *American Economic Review* 103 (4): 1507–28.
- Crispolti, Valerio, David Amaglobeli, and Xuguang Simon Sheng. 2022. “Cross-Country Evidence on the Revenue Impact of Tax Reforms.” *IMF Working Papers* 2022 (199): 1.
- Dabla-Norris, Era, and Frederico Lima. 2023. “Macroeconomic Effects of Tax Rate and Base Changes: Evidence from Fiscal Consolidations.” *European Economic Review* 153 (April): 104399.
- Das, Shuvam, Davide Furceri, Nikhil Patel, and Adrian Peralta-Alva. 2026. *AI Meets Fiscal Policy*. {SSRN} {Scholarly} {Paper}. Rochester, NY: Social Science Research Network.
- Dell, Melissa. 2025. “Deep Learning for Economists.” *Journal of Economic Literature* 63 (1): 5–58.
- Deng, Chunyuan, Yilun Zhao, Yuzhao Heng, et al. 2024. “Unveiling the Spectrum of Data Contamination in Language Model: A Survey from Detection to Remediation.” In *Findings of the Association for Computational Linguistics: ACL 2024*, edited by Lun-Wei Ku, Andre Martins, and Vivek Srikumar, 16078–92. Bangkok, Thailand: Association for Computational Linguistics.
- Dube, Arindrajit, Daniele Girardi, Òscar Jordà, and Alan M. Taylor. n.d. *A Local Projections Approach to Difference-in-Differences Event Studies*.
- Egami, Naoki, Musashi Hinck, Brandon Stewart, and Hanying Wei. 2023. “Using Imperfect Surrogates for Downstream Inference: Design-Based Supervised Learning for Social Science Applications of Large Language Models.” *Advances in Neural Information Processing Systems* 36 (December): 68589–601.
- Fellegi, Ivan P., and Alan B. Sunter. 1969. “A Theory for Record Linkage.” *Journal of the American Statistical Association* 64 (328): 1183–210.
- Fernandez-Fuertes, Ruben. 2025. *Monetary Policy Shocks: A New Hope — Large Language Models and Central Bank Communication*. {SSRN} {Scholarly} {Paper}. Rochester, NY: Social Science Research Network.

- Feuerriegel, Stefan, Abdurahman Maarouf, Dominik Bär, et al. 2025. "Using Natural Language Processing to Analyse Text Data in Behavioural Science." *Nat Rev Psychol* 4 (2): 96–111.
- Frankel, Jeffrey A., Carlos A. Végh, and Guillermo Vuletin. 2013. "On Graduation from Fiscal Procyclicality." *Journal of Development Economics* 100 (1): 32–47.
- Fritsch, Gabriel P., and J. Zachary Mazlish. n.d. *High-Frequency Fiscal Shocks*.
- Gambacorta, Leonardo, Byeungchun Kwon, Taejin Park, Pietro Patelli, and Sonya Zhu. 2024. *CB-LMs: Language Models for Central Banking*. nos. No 1215 (October).
- Gechert, Sebastian. 2015. "What Fiscal Policy Is Most Effective? A Meta-Regression Analysis." *Oxford Economic Papers* 67 (3): 553–80.
- Gentzkow, Matthew, Bryan Kelly, and Matt Taddy. 2019. "Text as Data." *Journal of Economic Literature* 57 (3): 535–74.
- Gilardi, Fabrizio, Meysam Alizadeh, and Maël Kubli. 2023. "ChatGPT Outperforms Crowd Workers for Text-Annotation Tasks." *Proceedings of the National Academy of Sciences* 120 (30): e2305016120.
- Golchin, Shahriar, and Mihai Surdeanu. 2024. "Time Travel in LLMs: Tracing Data Contamination in Large Language Models." *International Conference on Learning Representations* 2024 (May): 43008–29.
- Grimmer, Justin, Margaret E. Roberts, and Brandon M. Stewart. 2022. *Text as Data: A New Framework for Machine Learning and the Social Sciences*. Princeton University Press.
- Guajardo, Jaime, Daniel Leigh, and Andrea Pescatori. 2014. "Expansionary Austerity? International Evidence." *Journal of the European Economic Association* 12 (4): 949–68.
- Halterman, Andrew, and Katherine A. Keith. 2025. "Codebook LLMs: Evaluating LLMs as Measurement Tools for Political Science Concepts." *Political Analysis*, September, 1–17.
- Hebous, Shafik, and Tom Zimmermann. 2018. "Revisiting the Narrative Approach of Estimating Tax Multipliers." *Scandinavian J Economics* 120 (2): 428–39.
- Hsiang, Solomon M., Marshall Burke, and Edward Miguel. 2013. "Quantifying the Influence of Climate on Human Conflict." *Science* 341 (6151): 1235367.
- Ilzetzi, Ethan, Enrique G. Mendoza, and Carlos A. Végh. 2013. "How Big (Small?) Are Fiscal Multipliers?" *Journal of Monetary Economics* 60 (2): 239–54.
- Jha, Shikha, Sushanta K. Mallick, Donghyun Park, and Pilipinas F. Quising. 2014. "Effectiveness of Countercyclical Fiscal Policy: Evidence from Developing Asia." *Journal of Macroeconomics* 40 (June): 82–98.
- Jordà, Òscar. 2005. "Estimation and Inference of Impulse Responses by Local Projections." *American Economic Review* 95 (1): 161–82.

———. n.d. *Local Projections for Applied Economics*.

Kaminsky, Graciela L., Carmen M. Reinhart, and égh Carlos A. V. 2004. “When It Rains, It Pours: Procyclical Capital Flows and Macroeconomic Policies.” *NBER Macroeconomics Annual* 19 (January): 11–53.

Kim, Jeong-Hee. 2015. *Understanding Narrative Inquiry: The Crafting and Analysis of Stories as Research*. SAGE Publications.

Korinek, Anton. 2025. *AI Agents for Economic Research*. Working {Paper}. Working Paper Series. National Bureau of Economic Research.

Latifi, Albina, Viktoriia Naboka-Krell, Peter Tillmann, and Peter Winker. n.d. *Fiscal Policy in the Bundestag: Textual Analysis and Macroeconomic Effects*.

Ludwig, Jens, Sendhil Mullainathan, and Ashesh Rambachan. 2025. *Large Language Models: An Applied Econometric Framework*. arXiv.

Mertens, Karel, and Morten O. Ravn. 2013. “The Dynamic Effects of Personal and Corporate Income Tax Changes in the United States.” *American Economic Review* 103 (4): 1212–47.

———. 2014. “A Reconciliation of SVAR and Narrative Estimates of Tax Multipliers.” *Journal of Monetary Economics* 68 (December): S1–19.

Pescatori, Andrea, APescatori@imf.org, Daniel Leigh, et al. 2011. “A New Action-Based Dataset of Fiscal Consolidation.” *IMF Working Papers* 11 (128): 1.

Ramey, Valerie A. 2011. “Identifying Government Spending Shocks: It’s All in the Timing*.” *The Quarterly Journal of Economics* 126 (1): 1–50.

———. 2016. “Macroeconomic Shocks and Their Propagation.” In *Handbook of Macroeconomics*, 2:71–162. Elsevier.

———. 2019. “Ten Years After the Financial Crisis: What Have We Learned from the Renaissance in Fiscal Research?” *Journal of Economic Perspectives* 33 (2): 89–114.

Ramey, Valerie A., and Sarah Zubairy. 2018. “Government Spending Multipliers in Good Times and in Bad: Evidence from US Historical Data.” *Journal of Political Economy* 126 (2): 850–901.

Rathje, Steve, Dan-Mircea Mirea, Ilia Sucholutsky, Raja Marjeh, Claire E. Robertson, and Jay J. Van Bavel. 2024. “GPT Is an Effective Tool for Multilingual Psychological Text Analysis.” *Proceedings of the National Academy of Sciences* 121 (34): e2308950121.

Riera-Crichton, Daniel, Carlos A. Vegh, and Guillermo Vuletin. 2016. “Tax Multipliers: Pitfalls in Measurement and Identification.” *Journal of Monetary Economics* 79 (May): 30–48.

Romer, Christina D., and David H. Romer. 2017. “New Evidence on the Aftermath of Financial Crises in Advanced Countries.” *American Economic Review* 107 (10): 3072–118.

———. 2019a. *Fiscal Space and the Aftermath of Financial Crises: How It Matters and Why*. Working {Paper}. Working Paper Series. National Bureau of Economic Research.

———. 2019b. *Replication Data for: The Macroeconomic Effects of Tax Changes: Estimates Based on a New Measure of Fiscal Shocks*. Inter-university Consortium for Political; Social Research (ICPSR).

———. 2023. “Presidential Address: Does Monetary Policy Matter? The Narrative Approach After 35 Years.” *American Economic Review* 113 (6): 1395–423.

Romer, Christina D., and David H. Romer. 2010. “The Macroeconomic Effects of Tax Changes: Estimates Based on a New Measure of Fiscal Shocks.” *American Economic Review* 100 (3): 763–801.

———. n.d. *A Narrative Analysis of Postwar Tax Changes*.

Romer, Paul. n.d. *The Trouble With Macroeconomics*.

Singh, Harman, Nitish Gupta, Shikhar Bharadwaj, Dinesh Tewari, and Partha Talukdar. 2024. “IndicGenBench: A Multilingual Benchmark to Evaluate Generation Capabilities of LLMs on Indic Languages.” In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, edited by Lun-Wei Ku, Andre Martins, and Vivek Srikumar, 11047–73. Bangkok, Thailand: Association for Computational Linguistics.

Singh, Manasvini. 2026. *Deadly Stigma*. Working {Paper}. Working Paper Series. National Bureau of Economic Research.

Tillmann, Peter, Peter Winker, Albina Latifi, and Viktoriia Naboka-Krell. 2024. *Disagreement about Fiscal Policy* *. {SSRN} {Scholarly} {Paper}. Rochester, NY: Social Science Research Network.

Törnberg, Petter. 2025. “Large Language Models Outperform Expert Coders and Supervised Classifiers at Annotating Political Social Media Messages.” *Social Science Computer Review* 43 (6): 1181–95.

Torreblanca, Carolina, William Dinneen, Guy Grossman, and Yiqing Xu. 2025. *The Credibility Revolution in Political Science*. SocArXiv.

Vegh, Carlos A., and Guillermo Vuletin. 2015. “How Is Tax Policy Conducted over the Business Cycle?” *American Economic Journal: Economic Policy* 7 (3): 327–70.

Xuan, Weihao, Rui Yang, Heli Qi, et al. 2025. “MMLU-ProX: A Multilingual Benchmark for Advanced Large Language Model Evaluation.” In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, edited by Christos Christodoulopoulos, Tanmoy Chakraborty, Carolyn Rose, and Violet Peng, 1513–32. Suzhou, China: Association for Computational Linguistics.

Zhao, WeiKe, Chaoyi Wu, Yanjie Fan, et al. 2026. “An Agentic System for Rare Disease Diagnosis with Traceable Reasoning.” *Nature*, February, 1–10.

Ziems, Caleb, William Held, Omar Shaikh, Jiaao Chen, Zhehao Zhang, and Diyi Yang. 2024. “Can Large Language Models Transform Computational Social Science?” *Computational Linguistics* (Cambridge, MA) 50 (1): 237–91.