

C1: Measure Identification

H&K Stages S0-S3 — final codebook (v0.7.0) and development history

Esteban Degetau
World Bank
estebandegetau@gmail.com

Agustín Samano
World Bank
asamanope-
naloza@worldbank.org

June 28, 2026

C1 implements R&R Step 2 (Measure Identification): scanning a document chunk and deciding whether any portion of it substantively discusses a specific fiscal measure. The current deliverable is **v0.7.0**, accepted for deployment at S3 iteration 33. This notebook has two parts: the full final-iteration battery (mimicking the Halterman & Keith presentables) and the development history that produced it. The cross-codebook rollup lives in `iteration_summary.qmd`.

Part 1: The Final Codebook (v0.7.0)

S0: Codebook and Prompt

C1 is a binary classification task with two classes:

- **FISCAL_MEASURE**: any portion of the chunk substantively discusses a specific enacted, proposed, or under-consideration fiscal action.
- **NOT_FISCAL_MEASURE**: everything else (general economic commentary, aggregates, administrative procedure, bare name-drops).

Under v0.7.0 the positive class additionally returns a structured `measures[]` array (one entry per distinct act, with a per-measure country tag and three relevance flags), which is what downstream codebooks consume.

Table 1: C1 codebook structure (live `c1_codebook`).

Field	Value
Name	C1: Fiscal Measure Identification
Version	0.7.0
Classes	2
Clarifications per class	6, 4
Neg. clarifications per class	5, 2

Evaluation negatives

The zero-shot evaluation draws stratified hard negatives from `us_body`, designed to test the most common confusion patterns.

Table 2: Stratified hard-negative pool for evaluation.

Type	Count	%
economic_trends	90	30.0
extensions	60	20.0
proposals	120	40.0
withholding	30	10.0

Table 3: Evaluation data sources.

Source	Count
Positive (aligned_data passages)	368
Negative (stratified hard negatives)	300
Total	668

The full prompt

This is the assembled system prompt the model saw for the US evaluation run, with the `{country}` / `{country_iso}` tokens resolved (here, the United States). At deployment the same template is re-rendered per country.

i Prompt excerpt from C1

You are an expert in fiscal policy analysis classifying chunks from official government economic and budget documents. Your task is to scan the entire chunk and determine whether **any** portion of it substantively discusses a specific fiscal measure – enacted, proposed, or under consideration – that changes or would change fiscal liabilities or obligations. Chunks from government documents typically contain extended economic analysis, forecasts, and policy discussion – a chunk should be classified as FISCAL_MEASURE if it contains even one substantive discussion of a specific fiscal measure, regardless of what the rest of the chunk discusses. Read the class definitions carefully and classify the chunk using only the provided labels.

The chunk you are analysing comes from a corpus published by the government of **United States** (ISO code: `US`). Some chunks discuss fiscal measures enacted by other countries (foreign comparators, supranational bodies); these still qualify as FISCAL_MEASURE because they describe fiscal policy, but their country-of-enactment is `OTHER`, not `US`. The country-of-enactment tag is applied per measure, since a single chunk may discuss measures from multiple countries.

Class Definitions

FISCAL_MEASURE

****Definition:**** A chunk that contains, in any portion of its text, a substantive discussion of a specific legislative or executive action – whether enacted, proposed, or under consideration – that changes or would change fiscal liabilities or obligations.

****Inclusion criteria:****

- The passage describes a change in fiscal liabilities from taxes of any type – income taxes, payroll and social insurance taxes, consumption and excise taxes, investment incentives, and so on
- The action may be enacted into law, implemented by executive or administrative action (such as

S1: Behavioral Tests (H&K Figure 3)

S1 checks that the model emits legal labels, can recover the class definitions from memory, and is insensitive to class ordering. Test III (in-context example recovery) is **skipped** for C1: v0.7.0 carries no in-codebook examples.




Figure 1: S1 behavioral tests for C1 v0.7.0. Tests I/II shown as percent correct; Test IV shown as Fleiss' kappa with Landis (1977) reference bands.

Table 4: S1 results for the latest C1 iteration.

Test	Value	Threshold	Status
I_legal_outputs	100%	$\geq 100\%$	Pass
II_definition_recovery	100%	$\geq 100\%$	Pass
IV_order_invariance	0.0% change	$< 5\%$	Pass

S2: Zero-Shot Evaluation

Tier 1 chunks carry verbatim R&R passages (high-confidence labels); Tier 2 chunks are name-matched via the IKA heuristic (noisier). Tier 1 recall is the metric we trust most.

Table 5: S2 zero-shot metrics with 95% bootstrap CIs (latest C1 iteration).

Metric	Estimate	Target	Status
combined_recall	0.800 [0.771, 0.829]	≥ 0.90	Fail
tier1_recall	0.882 [0.803, 0.955]	≥ 0.95	Fail
precision	0.947 [0.929, 0.964]	≥ 0.70	Pass
tier2_recall	0.792 [0.762, 0.822]	—	—
f1	0.867 [0.849, 0.887]	—	—
accuracy	0.784 [0.757, 0.814]	—	—
specificity	0.670 [0.578, 0.761]	—	—

Confusion matrix

Table 6: S2 confusion matrix (live `c1_s2_eval`). Rows = predicted, columns = true.

Predicted	FISCAL_MEASURE	NOT_FISCAL_MEASURE
FISCAL_MEASURE	592	33
NOT_FISCAL_MEASURE	148	67

Per-act recall

Acts with passage-level recall below 1 may indicate codebook weaknesses or (more often at Tier 2) noisy labels.

```
c1_s2_eval$act_recall).
```

4 AMENDMENTS OF
 5 6 6 1980 1
 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25
 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47 48 49 50 51 52 53 54 55 56 57 58 59 60 61 62 63 64 65 66 67 68 69 70 71 72 73 74 75 76 77 78 79 80 81 82 83 84 85 86 87 88 89 90 91 92 93 94 95 96 97 98 99 100 101 102 103 104 105 106 107 108 109 110 111 112 113 114 115 116 117 118 119 120 121 122 123 124 125 126 127 128 129 130 131 132 133 134 135 136 137 138 139 140 141 142 143 144 145 146 147 148 149 150 151 152 153 154 155 156 157 158 159 160 161 162 163 164 165 166 167 168 169 170 171 172 173 174 175 176 177 178 179 180 181 182 183 184 185 186 187 188 189 190 191 192 193 194 195 196 197 198 199 200 201 202 203 204 205 206 207 208 209 210 211 212 213 214 215 216 217 218 219 220 221 222 223 224 225 226 227 228 229 230 231 232 233 234 235 236 237 238 239 240 241 242 243 244 245 246 247 248 249 250 251 252 253 254 255 256 257 258 259 260 261 262 263 264 265 266 267 268 269 270 271 272 273 274 275 276 277 278 279 280 281 282 283 284 285 286 287 288 289 290 291 292 293 294 295 296 297 298 299 300 301 302 303 304 305 306 307 308 309 310 311 312 313 314 315 316 317 318 319 320 321 322 323 324 325 326 327 328 329 330 331 332 333 334 335 336 337 338 339 340 341 342 343 344 345 346 347 348 349 350 351 352 353 354 355 356 357 358 359 360 361 362 363 364 365 366 367 368 369 370 371 372 373 374 375 376 377 378 379 380 381 382 383 384 385 386 387 388 389 390 391 392 393 394 395 396 397 398 399 400 401 402 403 404 405 406 407 408 409 410 411 412 413 414 415 416 417 418 419 420 421 422 423 424 425 426 427 428 429 430 431 432 433 434 435 436 437 438 439 440 441 442 443 444 445 446 447 448 449 450 451 452 453 454 455 456 457 458 459 460 461 462 463 464 465 466 467 468 469 470 471 472 473 474 475 476 477 478 479 480 481 482 483 484 485 486 487 488 489 490 491 492 493 494 495 496 497 498 499 500 501 502 503 504 505 506 507 508 509 510 511 512 513 514 515 516 517 518 519 520 521 522 523 524 525 526 527 528 529 530 531 532 533 534 535 536 537 538 539 540 541 542 543 544 545 546 547 548 549 550 551 552 553 554 555 556 557 558 559 560 561 562 563 564 565 566 567 568 569 570 571 572 573 574 575 576 577 578 579 580 581 582 583 584 585 586 587 588 589 590 591 592 593 594 595 596 597 598 599 600 601 602 603 604 605 606 607 608 609 610 611 612 613 614 615 616 617 618 619 620 621 622 623 624 625 626 627 628 629 630 631 632 633 634 635 636 637 638 639 640 641 642 643 644 645 646 647 648 649 650 651 652 653 654 655 656 657 658 659 660 661 662 663 664 665 666 667 668 669 670 671 672 673 674 675 676 677 678 679 680 681 682 683 684 685 686 687 688 689 690 691 692 693 694 695 696 697 698 699 700 701 702 703 704 705 706 707 708 709 710 711 712 713 714 715 716 717 718 719 720 721 722 723 724 725 726 727 728 729 730 731 732 733 734 735 736 737 738 739 740 741 742 743 744 745 746 747 748 749 750 751 752 753 754 755 756 757 758 759 760 761 762 763 764 765 766 767 768 769 770 771 772 773 774 775 776 777 778 779 780 781 782 783 784 785 786 787 788 789 790 791 792 793 794 795 796 797 798 799 800 801 802 803 804 805 806 807 808 809 810 811 812 813 814 815 816 817 818 819 820 821 822 823 824 825 826 827 828 829 830 831 832 833 834 835 836 837 838 839 840 841 842 843 844 845 846 847 848 849 850 851 852 853 854 855 856 857 858 859 860 861 862 863 864 865 866 867 868 869 870 871 872 873 874 875 876 877 878 879 880 881 882 883 884 885 886 887 888 889 890 891 892 893 894 895 896 897 898 899 900 901 902 903 904 905 906 907 908 909 910 911 912 913 914 915 916 917 918 919 920 921 922 923 924 925 926 927 928 929 930 931 932 933 934 935 936 937 938 939 940 941 942 943 944 945 946 947 948 949 950 951 952 953 954 955 956 957 958 959 960 961 962 963 964 965 966 967 968 969 970 971 972 973 974 975 976 977 978 979 980 981 982 983 984 985 986 987 988 989 990 991 992 993 994 995 996 997 998 999 1000 1001 1002 1003 1004 1005 1006 1007 1008 1009 1010 1011 1012 1013 1014 1015 1016 1017 1018 1019 1020 1021 1022 1023 1024 1025 1026 1027 1028 1029 1030 1031 1032 1033 1034 1035 10

Multi-measure act recall

A v0.7.0-specific secondary diagnostic: for chunks discussing multiple acts, did the model list each act?

Table 8: Per-act recall on multi-act chunks (live
c1_s2_eval\$multi_measure_act_recall).

Act	Chunks	Recalled	Recall
Changes in Depreciation Guide- lines and Revenue Act of 1962	2	0	0.00
Crude Oil Windfall Profit Tax Act of 1980	2	0	0.00
Excise Tax Reduction Act of 1954	2	0	0.00
Excise Tax Reduction Act of 1965	1	0	0.00
Expiration of Excess Profits Tax and of Temporary Income Tax In- creases	4	0	0.00
Federal-Aid Highway Act of 1956	1	0	0.00
Federal-Aid Highway Act of 1959	1	0	0.00
Jobs and Growth Tax Relief Rec- onciliation Act of 2003	1	0	0.00
Omnibus Budget Reconciliation Act of 1987	1	0	0.00
Public Law 89-800 (Suspension of Investment Tax Credit)	2	0	0.00
Public Law 90-26 (Restoration of the Investment Tax Credit)	3	0	0.00
Revenue Act of 1948	1	0	0.00
Revenue Act of 1964	4	0	0.00
Revenue and Expenditure Con- trol Act of 1968	2	0	0.00
Social Security Amendments of 1965	1	0	0.00
Taxpayer Relief Act of 1997 and Balanced Budget Act of 1997	1	0	0.00
Tax Reform Act of 1969	7	1	0.14
Revenue Act of 1945	5	1	0.20
Revenue Act of 1978	4	1	0.25
Tax Reduction Act of 1975	3	1	0.33

Error examples

Table 9: Sample of S2 misclassifications (live
c1_s2_eval\$error_analysis).

Act	Year	Type	True	Predicted
NA	1961	negativ	NOT_FISCAL_MEASURE	FISCAL_MEASURE
NA	1954	negativ	NOT_FISCAL_MEASURE	FISCAL_MEASURE
NA	1960	negativ	NOT_FISCAL_MEASURE	FISCAL_MEASURE
NA	2007	negativ	NOT_FISCAL_MEASURE	FISCAL_MEASURE
NA	2002	negativ	NOT_FISCAL_MEASURE	FISCAL_MEASURE
NA	1980	negativ	NOT_FISCAL_MEASURE	FISCAL_MEASURE
NA	1973	negativ	NOT_FISCAL_MEASURE	FISCAL_MEASURE
NA	1967	negativ	NOT_FISCAL_MEASURE	FISCAL_MEASURE
NA	1992	negativ	NOT_FISCAL_MEASURE	FISCAL_MEASURE
NA	1978	negativ	NOT_FISCAL_MEASURE	FISCAL_MEASURE
NA	1974	negativ	NOT_FISCAL_MEASURE	FISCAL_MEASURE
NA	1974	negativ	NOT_FISCAL_MEASURE	FISCAL_MEASURE
NA	1986	negativ	NOT_FISCAL_MEASURE	FISCAL_MEASURE
NA	2006	negativ	NOT_FISCAL_MEASURE	FISCAL_MEASURE
NA	1974	negativ	NOT_FISCAL_MEASURE	FISCAL_MEASURE

S3: Error Analysis (H&K Figure 4)

S3 perturbs the codebook to probe robustness: Test V (exclusion-criteria minimal pairs), Test VI (generic labels), Test VII (swapped label names).




Figure 2: S3 behavioral tests for C1 v0.7.0 (Tests V–VII), with the full-codebook baseline F1 as a reference line.

Ablation (H&K Table 4)

Components ranked by the accuracy/F1 drop when removed. Larger drops indicate more load-bearing components.

Table 10: Component ablation for the latest C1 iteration.

Condition	Accuracy	F1	Δ Accuracy	Δ F1
Full codebook	0.675	0.683	—	—
No label definitions	0.675	0.683	0.000	—
No clarifications	0.625	0.706	0.050	—
All removed	0.625	0.694	0.050	—

Multi-measure schema diagnostics

Three v0.7.0 schema checks (live [c1_s3_results](#)): are we over-listing measures, mis-tagging country of enactment, or under-listing on multi-act chunks?

Table 11: Over-listing: distribution of measures listed per FISCAL_MEASURE chunk.

Quantity	Value
Mean measures / chunk	4.50
Median (p50)	4
p90	7
Max	8
Chunks with ≥ 3 measures	6
FISCAL_MEASURE chunks	8

Table 12: Country-of-enactment tag distribution across listed measures.

Quantity	Value
Domestic (US) measures	89
OTHER (foreign) measures	0
Chunks containing any OTHER measure	0

Table 13: Under-listing: per-act recall on chunks expected to name multiple acts.

Quantity	Value
Multi-act chunks tested	3
Mean per-act recall	0.17
Chunks with full recall	0

Manual Error Analysis (H&K Table 5)

The actual S3 gate is a human read of all 40 baseline chunks against the H&K six-category framework (A–F). This is the headline result for v0.7.0.

Table 14: Manual error-analysis category distribution (latest C1 iteration).

Category	Count	Proportion
A: LLM correct	27	0.68
B: Incorrect gold standard	7	0.17
C: Document error	0	0.00
D: LLM non-compliance	0	0.00
E: Semantics/reasoning mistake	2	0.05
F: Other	4	0.10

After excluding Category B (incorrect gold standard), the label-noise-adjusted performance is:

Table 15: Bias-corrected gate metrics (Category B excluded).

Quantity	Value
Effective N	33
True positives	14
True negatives	13
False positives	4
False negatives	2
Accuracy	81.8%
Precision	77.8%
Recall	87.5%
Tier 1 recall	80.0%
Tier 2 recall	100.0%
Specificity	76.5%

The two Tier 1 false negatives in v0.7.0 (a Category E pair) are genuine semantic errors: the model recognized the relevant material but applied a stricter “substantive” bar than the codebook authorizes. The hypothesized cause is that the multi-measure schema raises the implicit threshold for what counts as a measure worth listing. The gaps were judged non-blocking and recorded for a future C1 revision; v0.7.0 was accepted for deployment.

Part 2: Development History

Everything below is sourced from `prompts/iterations/c1.yml` via the parsed `iteration_logs` target, so editing the log re-renders these figures and the narrative.

Timeline

Table 16: C1 iteration timeline (all runs, including exploratory non-Claude models).

Iter	Version	Stage	Date	Model	Pass
1	0.1.0	S1	2026-02-20	claude-haiku-4-5	✗
2	0.1.0	S1	2026-02-20	claude-haiku-4-5	✗
3	0.1.0	S1	2026-02-21	claude-haiku-4-5	✓
4	0.2.0	S1	2026-02-28	qwen/qwen-2.5-72b-instruct	✗
5	0.2.0	S1	2026-03-02	qwen/qwen-2.5-72b-instruct	✗
6	0.2.0	S2	2026-03-03	qwen/qwen-2.5-72b-instruct	✗
7	0.2.0	S2	2026-03-03	meta-llama/llama-3.3-70b-instruct	✗
8	0.2.0	S1	2026-03-03	claude-haiku-4-5	✓
9	0.2.0	S2	2026-03-03	claude-haiku-4-5	✗
10	0.3.0	S1	2026-03-03	claude-haiku-4-5	✓
11	0.3.0	S2	2026-03-03	claude-haiku-4-5	✗
12	0.4.0	S2	2026-03-04	claude-haiku-4-5	✗
13	0.4.0	S3	2026-03-09	meta-llama/llama-3.3-70b-instruct	—
14	0.4.0	S3	2026-03-09	meta-llama/llama-3.3-70b-instruct	—
15	0.4.0	S3	2026-03-09	claude-haiku-4-5	—
16	0.5.0	S3	2026-03-10	claude-haiku-4-5	—
17	0.5.0	S1	2026-04-04	claude-haiku-4-5	✓
18	0.5.0	S1	2026-04-05	claude-haiku-4-5	✗
19	0.5.0	S1	2026-04-05	gemma-2.5-flash	✗

Performance trajectory

The S2 zero-shot trajectory across formal (Claude) iterations. The decisive jump came with the v0.3.0 “detection reframe” (recall 29% → 81%), and the v0.4.0 relevance-filter reframe pushed Tier 1 recall to ceiling.




Figure 3: C1 Tier 1 recall across formal iterations.




Figure 4: C1 combined recall across formal iterations.




Figure 5: C1 precision across formal iterations.

What we learned, by version

For each codebook version, the concluding iteration’s interpretation and decision (verbatim from the log).

v0.1.0 — concluded at iteration 3 (S1, 2026-02-21)

What we found. C1 passed all S1 behavioral tests. The codebook example fix was the root cause of the Test III failure — the old negative example had reasoning that contradicted its structural label, and the model followed the reasoning. Test IV also now passes (0% change rate vs 10% before), likely because the tightened NE-1 reasoning removed the hedge language that made the model order-sensitive on the Budget Enforcement Act passage.

Decision. Proceed to S2 (LOOCV evaluation).

v0.2.0 — concluded at iteration 9 (S2, 2026-03-03)

What we found. First formal S2 zero-shot evaluation with Haiku on codebook v0.2.0. Precision is near-perfect (99.1%) but recall is catastrophically low (29.1%, target $\geq 90\%$). The model applies a “dominant topic” heuristic rather than scanning for fiscal measures within chunks. 541/543 errors are false negatives with consistent reasoning patterns: “economic analysis” (27%), “general economic” (23%), “does not describe” (21%), “projections” (18%). Even Tier 1 chunks (verbatim passage matches) achieve only 42.9% recall – the model acknowledges fiscal measures exist in the text but dismisses them because the chunk’s overall character is economic analysis. 120/238 act-year rows have 0% recall (114 are Tier 2 only). The 2 false positives appear to be mislabeled negatives (Treasury Annual Reports containing genuine fiscal measures).

Root cause: H&K Error Category B (prompt construction) – the codebook’s “substantive detail” requirement and negative clarifications give the model license to dismiss chunks where the fiscal measure is a minority of the content. Formal S3 skipped for this iteration because the failure mode is unambiguous and ablation would rank components of a codebook being retired.

Decision. Revise codebook to v0.3.0 with two targeted changes, then re-run S1 and S2. (1) Reframe instructions from topic classification to detection. Change the instructions to explicitly say: “Scan the entire chunk for any portion that substantively describes an enacted fiscal measure. A chunk should be classified as FISCAL_MEASURE even if the majority of the text discusses other economic topics – one substantive description anywhere in the chunk is sufficient.” (2) Soften the “substantive detail” bar but keep it meaningful. Instead of requiring “specific rate changes, affected populations, revenue estimates” (which the model interprets as a very high bar), say something like: “The passage names a specific enacted act or executive action and provides enough context to understand what fiscal liabilities it changed – this could include rate changes, affected populations, revenue estimates, effective dates, or even a general description of the measure’s purpose and scope.” This lets Tier 2 chunks (which name the act and describe it generally) pass without collapsing the bar to mere mentions. Run S3 only if v0.3.0 S2 results are in the right ballpark but need refinement.

v0.3.0 – concluded at iteration 11 (S2, 2026-03-03)

What we found. v0.3.0 detection reframe massively improved recall (+52 pp combined, +40 pp Tier 1) vs v0.2.0 (iteration 9) while keeping precision high (95.1%). The dominant-topic heuristic that caused catastrophic recall in v0.2.0 is largely resolved. Remaining 144 FN are dominated by “not enacted/proposed” reasoning (87% of FN cite this) – the model over-applies the negative clarification about unenacted proposals to chunks where enacted acts are discussed in forward-looking or conditional tense. 32 FP are largely H&K Error Category F (ambiguous ground truth): genuine fiscal measures not in the 44-act label set, meaning reported precision is a conservative lower bound. 19 zero-recall act-year rows are all Tier 2-only single-chunk distant references. Tier 1 and Tier 2 miss rates are similar (~17-19%), suggesting the problem is the model’s enacted-status threshold, not reference vagueness. 0 parse failures (100% valid JSON from Haiku).

Decision. Revise codebook to v0.4.0 with a strategic reframe based on discussion of results. Three changes: (1) Remove the enacted/proposed filter from C1 entirely — C1 becomes a recall-optimized relevance filter for fiscal policy discussion; enacted status determination moves downstream to a future filter or C2. If we ex post find more acts than R&R, we add an enacted filter downstream. (2) Further soften the substantive detail bar to push Tier 1 recall toward 95%+; Tier 1 recall is the top priority because these chunks carry the richest information for C2-C4 classification. (3) Accept that reported precision is conservative due to Category F FPs in the negative test set. Precision cost is bearable because downstream codebooks (C2-C4) will filter irrelevant chunks. Additionally, improve the act identification pipeline to search for acronym mentions of known acts (e.g., ERTA, OBRA, TRA), improving label coverage for evaluation.

v0.4.0 — concluded at iteration 16 (S3_MANUAL_ANALYSIS, 2026-03-10)

What we found. There may be room for codebook improvement. The 3 Category E errors share a common pattern: the model conflates government financial actions (intergovernmental transfers, education appropriations, debt management) with taxpayer-liability-changing fiscal measures. The Category F case (EIA) reveals R&R boundary ambiguity on organizational acts with fiscal implications. The 5 Category B errors confirm that all false negatives are gold-standard noise from IKA Tier 2 matching, not codebook or model failures. Bias-corrected recall is 100%.

Decision. Consider improving the codebook with a clarification distinguishing changes to taxpayers’ liabilities from other government financial actions (intergovernmental transfers, debt management, organizational acts). This is the single highest-leverage codebook improvement identified by manual analysis.

v0.5.0 — concluded at iteration 20 (S3_MANUAL_ANALYSIS, 2026-04-04)

What we found. C1 v0.5.0 achieves zero semantic/reasoning errors (Category E = 0), up from 3 in v0.4.0. The taxpayer-liability clarification added in v0.5.0 fixed text 35 (TIPS). Bias-corrected recall is 100% across both tiers. All residual errors are Category F (codebook scope ambiguity): 3 false positives where the model correctly identified real fiscal content that falls outside R&R’s tax/revenue scope. These share a common root cause — C1’s definition of ‘fiscal measure’ does not encode R&R’s implicit boundary

between tax legislation and spending-side policy. A deeper insight: some of these measures live in a gray zone where spending policy is implemented through tax-adjacent mechanisms (credits, rebates), and their exclusion from R&R may be driven by motivation-classification needs in C2 rather than C1 scope alone. This is a potential source of C1 improvement, either through codebook clarification or pipeline design (downstream C2 filter). With $n=3$, this is a hypothesis to monitor during C2 development, not a confirmed systematic pattern.

Decision. C1 S3 gate passed. Proceed to C2 (Motivation) development. During C2 development, monitor the tax-vs-spending gray zone: observe how measures with behavioral/political motivation (energy policy, welfare reform) interact with R&R’s 4-bucket classification. This will inform whether the F-error pattern requires a C1 codebook fix or is better addressed as a C2-level filter.

v0.6.0 — concluded at iteration 28 (S3_MANUAL_ANALYSIS, 2026-04-07)

What we found. C1 v0.6.0 S3 regression check complete. Zero semantic errors ($E=0$), unchanged from v0.5.0. The single new prediction (text_id 24) is a positive finding: the extra_output_fields prompted the model to read footnotes more carefully, correctly identifying two R&R-relevant fiscal measures (SS Amendments 1950, Revenue Act 1951) that the v0.5.0 model and iteration 20’s human reviewer both overlooked. Bias-corrected recall is 100% across both tiers. The 3 residual F errors (texts 23, 32, 33) remain the same spending-side scope ambiguity from v0.5.0 — these are deferred to C2 development as planned in iteration 20.

Decision. C1 S3 gate passed for v0.6.0. Proceed to C2 (Motivation) codebook development. The extra_output_fields are a net positive — they improve model thoroughness without introducing new errors. During C2 development, monitor the 3 F-error pattern (spending-side fiscal policy) as planned.

v0.7.0 — concluded at iteration 33 (S3_MANUAL_ANALYSIS, 2026-05-20)

What we found. New schema either changes the model’s classification framework or distracts it from reasoning the main codebook task (fiscal shock identification). The 2 new Tier 1 E errors (text_ids 3, 7) are both cases where the model recognized the relevant content and then applied an overly restrictive bar — consistent with the hypothesis that asking for per-measure structure (name

+ country + 3 relevance flags) raises the model's threshold for what counts as a measure worth listing. Tier 2 recall remains at ceiling (100%) after excluding IKA noise; specificity drops modestly because the codebook-gap F-category continues to flag substantive foreign-credit and spending-side content as FISCAL_MEASURE. Net: v0.7.0 is fit-for-purpose for Phase 2 deployment despite the Tier 1 regression — the multi-measure schema unlocks the per-measure structure required for C2a (Phase 5) and Malaysia (Phase 2) downstream consumption.

Decision. Defer identified gap fixes to a future C1 version. Accept v0.7.0 as-is for deployment. Three gaps recorded for future revision (foreign-credit/capital-controls, spending-side authorizations, overly-restrictive substantive threshold) but not blocking. C1 S3 manual audit closes Step 4 of the 2026-05-18 sequencing plan in docs/deltas.md. Next: Step 5 (C2a v0.6.0).